

Risk-Controlled Spatio-Temporal Graph Learning for Anti-Money Laundering Under Extreme Imbalance and Topological Camouflage

Md. Nazmul, Musrat Jahan Gungun, and Maheli Ahmed

Abstract—Financial money laundering moves an estimated \$800B–\$2T annually across global financial rails. While Graph Neural Networks (GNNs) capture relational topologies, practical Anti-Money Laundering (AML) surveillance is hindered by velocity evasion, topological camouflage, extreme class imbalance ($\pi < 0.05\%$), cold-start isolation, and uncalibrated uncertainty. We propose C-STGB (Conformal Spatio-Temporal GraphBoost), an integrated risk-controlled temporal graph learning framework combining continuous harmonic temporal attention, learnable edge-trust gating, typology-clustered latent GraphSMOTE, evidence-adaptive graph-tabular fusion, and Class-Conditional Conformal Risk Control (CRC). Benchmarked across 14 financial transaction networks (9.53M entities, 9.53M transactions, 182 evaluation pairs over five locked seeds), C-STGB achieves an unweighted macro F1 of 67.26% (0.6848 PR-AUC). This establishes parity with tuned gradient-boosted decision trees (XGBoost 67.92%, $p = 0.8077$; CatBoost 67.14%, $p = 0.4318$) which excel on flat transaction streams, alongside statistically significant outperformance over evaluated deep GNNs (+43.49 pp mean uplift over GCN; $W = 105.0$, $p_{\text{adj}} < 0.001$), with advantages concentrated in graph-rich relational regimes. Furthermore, C-STGB provides finite-sample coverage ($\geq 99.0\%$) under within-class exchangeability with Adaptive Conformal Inference (ACI) drift tracking, routing $> 99.4\%$ of volume into decisive singleton prediction sets while achieving a streaming Fast-Path inference latency of 0.45 ms per event.

Index Terms—Data Mining, Graph Neural Networks, Temporal Graph Learning, Class Imbalance, Conformal Risk Control, Anomaly Detection, Anti-Money Laundering.

I. INTRODUCTION

FINANCIAL crime surveillance is a critical application of machine learning in global transaction infrastructure. According to the United Nations Office on Drugs and Crime (UNODC) and Financial Action Task Force (FATF), illicit networks launder between \$800 billion and \$2 trillion annually (2% to 5% of global GDP) across inter-bank wire rails, mobile financial service (MFS) e-wallets, and crypto-asset ledgers [1], [2]. Traditional rule-based Financial Intelligence Units (FIUs) trigger excessive false positive alerts ($>90\%$ to 95%), placing

an overwhelming operational burden on human compliance teams.

While Graph Neural Networks (GNNs) capture relational transaction topologies [3]–[7], practical AML surveillance encounters five central technical challenges: (1) **Continuous-Time Velocity Evasion**: microsecond transaction bursts alongside multi-week dormant holding chains cannot be captured simultaneously by discrete snapshot windowing (e.g., EvolveGCN [8]) or single exponential decay kernels; (2) **Adversarial Topological Camouflage**: illicit entities deliberately routing synthetic transactional chaff to high-degree commercial hubs induce severe feature over-smoothing during neighborhood message aggregation [9]; (3) **Extreme Class Imbalance** ($\pi < 0.05\%$): severe class skew causes posterior attenuation, causing minority recall collapse under default decision thresholds ($\tau = 0.50$); (4) **Cold-Start Graph Isolation** ($d(u) \leq 2$): newly activated accounts receive uninformative message vectors under spatial aggregation; and (5) **Uncalibrated Decision Uncertainty**: black-box scores lack statistical guarantees under non-stationary temporal distribution drift.

To address these challenges, we formulate and investigate three core hypotheses: (H_1) continuous multi-scale temporal attention coupled with learnable edge-trust filtering suppresses camouflage-induced over-smoothing across burst and dormancy dynamics, preserving minority discrimination relative to static and single-decay dynamic GNNs; (H_2) interpolating minority representations within latent typology clusters (C_k) recovers minority recall without elevating false-alarm rates on benign commercial hubs under extreme class skew; and (H_3) Class-Conditional Conformal Risk Control (CRC) delivers finite-sample label coverage ($\geq 99\%$) under within-class exchangeability with bounded drift degradation, routing the vast majority ($> 99\%$) of transaction volume into decisive singleton prediction sets.

Importantly, our empirical findings reveal that graph neural architectures do not universally surpass strong gradient-boosted decision tree baselines across all transaction streams. Rather, their principal advantage is concentrated in graph-rich, multi-hop relational crime settings, whereas optimized tree ensembles remain competitive or superior on predominantly flat, single-hop transactional data. Rebuilding AML surveillance around this empirical reality, we present **C-STGB** (Conformal Spatio-Temporal GraphBoost), organized around four primary scientific pillars:

- 1) **Continuous-Time Relational Representation**: Integrates continuous harmonic temporal encodings (Φ_{time}) with

Md. Nazmul (ORCID: 0009-0001-6115-7023) and Musrat Jahan Gungun (ORCID: 0009-0006-4249-9198) are with the Department of Computer Science and Engineering, College of Technology, National University, Gazipur 1704, Bangladesh (e-mail: nazmulhas36@gmail.com, gungun-jahan84@gmail.com).

Maheli Ahmed (Faculty Supervisor, ORCID: 0000-0002-5183-7498) is with the Department of Computer Science and Engineering, College of Technology, National University, Gazipur 1704, Bangladesh (e-mail: maheli.ahmed.cse.cot@gmail.com).

Corresponding author: Md. Nazmul.

Hawkes arrival velocity features and multi-scale attention, distinguishing microsecond bursts from dormant holding chains.

- 2) **Topology-Aware Robustness Under Camouflage:** Employs a context-aware learnable edge-trust gate ($\hat{g}_{ij}$) that attenuates spurious camouflage edges during message aggregation, mitigating feature over-smoothing around high-degree hubs.
- 3) **Typology-Aware Imbalance Learning:** Constrains latent topological oversampling along empirical laundering clusters (C_k , $K = 4$) with bilinear link generation and Asymmetric Focal Tversky loss, preserving point-level flow-balance regularity while boosting minority recall.
- 4) **Risk-Controlled Selective Prediction:** Formulates class-conditional conformal prediction for continuous transaction networks, establishing finite-sample coverage ($\geq 99.0\%$) under within-class exchangeability, reinforced by non-asymptotic drift bounds and streaming Adaptive Conformal Inference (ACI).

Empirical Findings & Paper Organization: Benchmarked across 14 financial transaction networks (9.53M entities, 9.53M transactions across 182 evaluation pairs over five locked seeds), C-STGB achieves 67.26% unweighted macro F1 (0.6848 PR-AUC), establishing statistically significant superiority over deep GNNs (+43.49 pp mean uplift over GCN, $W = 105.0$, $p_{\text{adj}} < 0.001$) and parity with tuned tree ensembles (CatBoost 67.14%, $p = 0.4318$; XGBoost 67.92%, $p = 0.8077$), with streaming Fast-Path inference operating at 0.45 ms per single event. The remainder of this paper is organized as follows: Section II reviews related work; Section III establishes the problem formulation and C-STGB architecture; Section IV presents empirical evaluations and controlled ablations; Section V discusses operational implications and domain boundaries; Section VI details forensic case studies and pilot compliance decision support; Section VII outlines limitations; and Section VIII concludes.

II. RELATED WORK

A. Graph Representation Learning for Anti-Money Laundering & Fraud

The application of Graph Neural Networks (GNNs) to financial forensic analysis was established by Weber *et al.* [7] on the Elliptic Bitcoin dataset, demonstrating that spatial architectures like Graph Convolutional Networks (GCN) [5] and GraphSAGE [6] capture localized neighborhood connectivity beyond isolated tabular classifiers. Subsequent extensions by Bellei *et al.* [10] introduced the Elliptic2 multi-asset dataset, emphasizing higher-order relational subgraphs, while Johannessen and Jullum [11] explored heterogeneous graph representations across commercial account ledgers. In card and banking networks, Cheng *et al.* [12] pioneered group-aware deep graph learning in IEEE TKDE to capture collective community-level transactions, while Fu *et al.* [13] introduced H²IDE in IEEE TKDE to untangle fraudulent camouflage on multi-relational graphs via disentangled homophily and heterophily identification. More broadly, Qiao *et al.* [14] surveyed modern deep graph anomaly detection paradigms in IEEE TKDE,

emphasizing the persistent challenges of structural camouflage and non-uniform relation topologies. Beyond UTXO graphs, graph representations have been applied to Ethereum smart contract security [15], [16], exchange insolvency forensics [17], multi-bank synthetic wire simulations [18], [19], streaming credit card fraud [20], and cross-institutional federated learning architectures [21]. However, prior AML graph learning frameworks largely operate over static snapshots or similarity graphs, remaining blind to continuous multi-scale temporal burstiness and susceptible to severe over-smoothing when launderers attach chaff edges to high-degree commercial hubs [9]. A comprehensive comparison across all evaluated model paradigms is detailed in Supplementary Table S2.

B. Continuous-Time Dynamic Temporal Graphs & Point Processes

Financial transaction networks evolve through continuous-time asynchronous events. Discrete snapshot dynamic GNNs, such as EvolveGCN [8], discretize event streams into regular time slices, discarding intra-slice temporal ordering and obscuring sub-second velocity bursts. While continuous frameworks including Temporal Graph Networks (TGN) [22], TGAT [23], and continuous anomaly models [24] preserve event order, standard single exponential decay kernels ($\exp(-\lambda\Delta t)$) rapidly attenuate to zero over multi-week dormant holding periods. Self-exciting point processes, formalized by Hawkes [25], effectively capture transaction clustering and burst velocity acceleration, while continual graph architectures counteract streaming non-stationarity [26]. Standard continuous-time GNNs employ uniform decay schedules that struggle to simultaneously represent sub-second transaction bursts alongside multi-week holding intervals characteristic of layering chains.

C. Adversarial Graph Learning & Topological Camouflage Defense

Adversarial attacks on graph neural networks manipulate connectivity or node attributes to induce misclassification [27], [28]. In financial forensics, adversaries employ topological camouflage: laundering entities deliberately attach transactional edges to high-degree, reputable commercial hubs (e.g., exchanges, payment gateways, utility merchants) to blend into background traffic [9]. Because standard graph convolutional aggregation treats incident edges with uniform or degree-normalized weights, camouflage connections drag illicit node embeddings toward majority benign centroids. CARE-GNN [9] introduced reinforcement learning to select informative neighbors on multi-relational graphs, while Ou *et al.* [29] recently developed CamFD to detect temporal behavioral inconsistencies of camouflaged fraudsters on dynamic graphs. However, existing camouflage defenses either rely on static relation heuristics or assume homophilic user interactions, lacking continuous-time soft-gating that dynamically prunes high-degree banking chaff while providing provable finite-sample risk coverage.

D. Imbalanced Graph Representation Learning & Latent Oversampling

Real-world financial transaction graphs exhibit extreme class imbalance, often with illicit prevalence $\pi < 0.05\%$. Traditional Euclidean oversampling techniques such as SMOTE [30] operate strictly on tabular feature vectors, generating synthetic instances without topologically valid relational edges. Under extreme skew, Precision-Recall (PR-AUC) analysis is necessary to prevent false optimism from dominant true negatives [31]. To bridge topological oversampling, Zhao *et al.* [32] proposed GraphSMOTE, performing feature interpolation in the intermediate latent embedding space and decoding synthetic edges via an auxiliary link generator. However, standard GraphSMOTE executes unconstrained latent interpolation across the global minority class pool, neglecting distinct structural laundering typologies (cyclic wash trading vs. smurfing fan-out vs. pass-through layering), risking synthetic instances that violate domain-level flow-balance regularity.

E. Conformal Risk Control, Selective Classification & Temporal Shift

Conformal prediction, formalized by Vovk *et al.* [33] and extended by Angelopoulos and Bates [34], [35], provides distribution-free, finite-sample uncertainty quantification. Recent research has investigated conformal prediction on GNNs [36] and class-conditional conformal prediction under class imbalance [37], [38]. While standard marginal conformal prediction guarantees valid overall coverage $(1 - \alpha)$, it frequently suffers from severe conditional under-coverage on the rare class in highly imbalanced datasets. Under temporal non-stationarity, standard exchangeability is violated and coverage must be explicitly bounded through distribution drift metrics [39]–[41].

F. Recent Uncertainty-Aware & Dynamic Financial Graph Learning

Recent comprehensive surveys by Zhang *et al.* [42] and Pourhabibi *et al.* [43] review over 100 financial graph fraud studies, establishing that enterprise AML systems must simultaneously address continuous temporal dynamics, adversarial camouflage, and reliable uncertainty quantification. Most recently, Chen *et al.* [44] proposed an uncertainty-aware spatio-temporal graph neural architecture combining continuous-time encoding with evidential deep learning for fraud triage. While evidential neural networks parameterize Dirichlet priors to approximate epistemic uncertainty, they rely on heuristic distributional assumptions and offer no finite-sample coverage guarantees. Unlike heuristic or evidential uncertainty frameworks [44], C-STGB leverages distribution-free Conformal Risk Control to deliver rigorous finite-sample class-conditional coverage bounds ($\geq 1 - \alpha$), backed by Martingale convergence under Adaptive Conformal Inference (ACI) tracking, while delineating the boundary where GNN message passing is indispensable vs. where tabular trees remain competitive (Supplementary Table S2).

III. PROBLEM FORMULATION & C-STGB METHODOLOGY

A. Problem Formulation and Adversarial Threat Model

We formalize financial crime surveillance over a continuous-time dynamic transaction graph $\mathcal{G}_t = (\mathcal{V}_t, \mathcal{E}_t, \mathbf{X}_v^{(t)}, \mathbf{X}_e^{(t)})$, where $\mathcal{V}_t$ represents active transacting accounts, $\mathcal{E}_t = \{(u, v, t_e, \mathbf{e}_{uv}) : t_e \leq t\}$ denotes directed fund transfers occurring at or before timestamp t , $\mathbf{X}_v^{(t)} \in \mathbb{R}^{|\mathcal{V}_t| \times d_v}$ comprises entity features, and $\mathbf{X}_e^{(t)} \in \mathbb{R}^{|\mathcal{E}_t| \times d_e}$ denotes edge features (transfer amount, currency rail, transaction fees). Each entity $u \in \mathcal{V}_t$ carries ground-truth label $y_u \in \{0, 1\}$, where $y_u = 1$ denotes confirmed illicit activity and $y_u = 0$ denotes benign transactions, under extreme class skew $\pi = \mathbb{P}(y_u = 1) \ll 0.01$. In transaction streams lacking explicit account graphs (e.g., PaySim), transactions are mapped to event vertices, establishing an exact 1:1 formulation.

Adversarial Threat Model: Adversarial syndicates evade forensic detection through two coordinated strategies: (1) *Velocity Obfuscation*: interweaving sub-minute structuring bursts with multi-week dormant holding periods to evade static temporal windows; and (2) *Topological Camouflage*: actively attaching transactional chaff to high-degree, reputable commercial hubs (e.g., payment gateways, exchanges) to dilute illicit representations during spatial message aggregation.

Operational Objectives: The surveillance engine optimizes two complementary objectives:

- 1) **Cost-Sensitive Classification:** Learn an inductive hypothesis $f_\Theta : \mathcal{V}_t \times \mathcal{G}_t \rightarrow [0, 1]$ predicting anomaly posterior $\hat{p}_u = f_\Theta(u | \mathcal{G}_t)$ that minimizes asymmetric regulatory Bayes risk.
- 2) **Risk-Controlled Queue Triage:** Construct a class-conditional conformal mapping $\Gamma : \mathcal{V}_t \rightarrow 2^{\{0,1\}}$ routing transactions to operational action sets $\Gamma(u) \subseteq \{0, 1\}$. Under within-class exchangeability, Γ provides finite-sample coverage $\mathbb{P}(y_u \in \Gamma(u) | y_u = y) \geq 1 - \alpha_y$ at specified significance $\alpha_y \in (0, 1)$, with provable bounds under temporal distribution shift. Decisive singletons are routed to automated workflows ($\Gamma(u) = \{1\}$ to Tier 1 Escalation; $\Gamma(u) = \{0\}$ to Tier 3 Straight-Through Clearance), restricting human officer investigation to ambiguous sets ($\Gamma(u) = \{0, 1\}$ in Tier 2).

All evaluated benchmark datasets operate under strict governance protocols: public cryptocurrency ledgers and exchange insolvency logs are indexed via cryptographic hashes without non-public personal information (PII), and synthetic benchmarks simulate multi-agent banking rails without proprietary customer accounts (detailed in Supplementary Section A). Mathematical symbols and dimensions are summarized in Supplementary Table S1.

B. 12-Dimensional Canonical Mass-Flow and Velocity Invariants

Design Rationale: Standard GNNs evaluate node features without domain mass-conservation constraints. In money laundering, intermediate pass-through mules require capital to enter and exit rapidly with near-zero retention. C-STGB formalizes 12-dimensional domain-informed mass-flow and relational invariants $\mathbf{z}_{\text{inv}}(u, t) \in \mathbb{R}^{12}$:




Fig. 1. Comprehensive 6-module architecture of C-STGB: (1) Continuous-Time Multi-Scale Temporal Graph Encoder; (2) Adaptive Edge Trust Filter; (3) Typology-Aware Minority Graph Augmentation with Hard-Negative Mining; (4) Evidence-Adaptive Graph-Tabular Fusion for Cold-Start Resolution; (5) Cost-Sensitive Risk Head; and (6) Class-Conditional Conformal Risk Control with Selective Abstention Triage.

1) Flow Regularity and Degree Asymmetry Ratios:

$$\Phi_{\text{flow}}(u, t) = \frac{\sum_{e \in \mathcal{E}_{\text{out}}(u)} \text{Amount}(e)}{\sum_{e \in \mathcal{E}_{\text{in}}(u)} \text{Amount}(e) + \epsilon}, \quad (1)$$

$$\Psi_{\text{asym}}(u, t) = \frac{|d_{\text{in}}(u) - d_{\text{out}}(u)|}{d_{\text{in}}(u) + d_{\text{out}}(u) + \epsilon} \quad (2)$$

with $\epsilon = 10^{-5}$. Intermediate mules operate near unity ($\Phi_{\text{flow}} \approx 1.0$), whereas legitimate commercial accounts exhibit accumulation or disbursement variance. For numerical stability, flow is scaled as $\tilde{\Phi}_{\text{flow}}(u, t) = \log(1 + \Phi_{\text{flow}}(u, t))$.

2) Hawkes Point Process Arrival Intensity:

$$\lambda_u(t) = \mu_u + \sum_{t_i < t} \alpha_{\text{hww}} \exp\left(-\beta_{\text{hww}} \frac{t - t_i}{24.0}\right) \quad (3)$$

parameterized over causal event history $[t - \Delta T, t]$ to model transaction burst acceleration. We enforce sub-critical branching $\alpha_{\text{hww}}/\beta_{\text{hww}} < 1.0$ and clamp $\tilde{\lambda}_u(t) = \log \min(\max(\lambda_u(t), 10^{-5}), 10^4)$.

3) Causal Historical Taint Propagation (Zero-Leakage):

$$\mathbf{s}_{\text{fwd}} = (1 - \alpha_{\text{ppr}})(\mathbf{I} - \alpha_{\text{ppr}}\mathbf{P}^T)^{-1}\mathbf{s}_{\text{seed}}, \quad (4)$$

$$\mathbf{s}_{\text{bwd}} = (1 - \alpha_{\text{ppr}})(\mathbf{I} - \alpha_{\text{ppr}}\mathbf{P})^{-1}\mathbf{s}_{\text{seed}} \quad (5)$$

computed exclusively over causal historical subgraphs ($t_e \leq t$). Seed indicator $\mathbf{s}_{\text{seed}}(u, t) = 1.0$ strictly if $u \in \mathcal{V}_{\text{illicit}}^{\text{train}}$ and confirmation timestamp $t_{\text{conf}}(u) \leq t - \Delta T_{\text{delay}}$, and 0.0 otherwise, where $\Delta T_{\text{delay}} \in [30, 90]$ days models investigative delay. Truncated Power Iteration ($T_{\text{iter}} = 3$, tolerance $\epsilon = 10^{-4}$) over degree-capped graph $\mathcal{G}_t^{(K)}$ ($K \leq 15$) ensures deterministic $O(K \cdot T_{\text{iter}})$ computation with < 0.08 ms overhead. All validation, calibration, and test nodes receive $\mathbf{s}_{\text{seed}}(u) \equiv 0.0$, strictly precluding label leakage.

4) Complete 12-D Descriptor Vector:

$$\mathbf{z}_{\text{inv}}(u, t) = [\tilde{\Phi}_{\text{flow}}, \Psi_{\text{asym}}, \lambda_u, \mathbf{s}_{\text{fwd}}, \mathbf{s}_{\text{bwd}}, \mu_A, \sigma_A^2, \gamma_A, \kappa_A, v_A, v_{\text{io}}, \sigma_{\Delta t}]^T \quad (6)$$

where $\{\mu_A, \sigma_A^2, \gamma_A, \kappa_A, v_A\}$ denote the 5-moment causal ego-statistics of transaction amounts, v_{io} is the input/output velocity ratio, and $\sigma_{\Delta t}$ is the inter-arrival temporal dispersion (Supplementary Fig. S5).

C. Continuous-Time Tri-Band Harmonic Temporal Attention

Design Rationale: Standard temporal GNNs use single exponential decay kernels ($\exp(-\lambda \Delta t)$), which either extinguish

rapidly over multi-week holding chains or blur sub-minute structuring bursts. C-STGB projects continuous time deltas $\Delta t = t_v - t_u \geq 0$ into a continuous harmonic positional space $\Phi_{\text{time}}(\Delta t) \in \mathbb{R}^{d_t}$ with $\omega_k = 10000^{-2k/d_t}$:

$$\Phi_{2k}(\Delta t) = \sin(\omega_k \log(1 + \Delta t)), \quad (7)$$

$$\Phi_{2k+1}(\Delta t) = \cos(\omega_k \log(1 + \Delta t)) \quad (8)$$

Dynamic temporal attention weights synthesize three distinct velocity bands:

$$w(\Delta t) = \sum_{b \in \mathcal{B}} \pi_b [w_{\text{min}} + (1 - w_{\text{min}})e^{-\lambda_b \Delta t}] \times (1 + \tanh(\gamma_b \lambda_u(t))) \quad (9)$$

where $\mathcal{B} = \{\text{burst, diurnal, seasonal}\}$ with learnable mixture weights $\sum_{b \in \mathcal{B}} \pi_b = 1.0$ via softmax. The burst band ($\lambda_{\text{burst}} = 0.80$) captures micro-structuring within minutes, the diurnal band ($\lambda_{\text{diurnal}} = 0.05$) models 24-hour business cycles, and the seasonal band ($\lambda_{\text{seasonal}} = 0.01$) maintains long-range memory of dormant holding wallets over 90 to 180 days.

D. Context-Aware Adaptive Edge Trust Filter & Unified Attention

Design Rationale: Adversarial launderers deliberately inject high-volume transactional chaff connecting to reputable commercial hubs. Because standard isotropic convolutions average neighbor representations, these camouflage edges drag illicit embeddings toward benign centroids. C-STGB computes a context-aware edge-trust gate $\hat{g}_{ij} \in [\delta_{\text{floor}}, 1.0]$:

$$\mathbf{c}_{ij} = [\mathbf{q}_i \| \mathbf{k}_j \| \mathbf{e}_{ij} \| \Phi_{\text{time}}(\Delta t) \| \mathbf{z}_{\text{inv}}(j, t)], \quad (10)$$

$$\hat{g}_{ij} = \delta_{\text{floor}} + (1 - \delta_{\text{floor}})\sigma\left(\mathbf{W}_{g2} \text{LeakyReLU}(\mathbf{W}_{g1}\mathbf{c}_{ij} + \mathbf{b}_{g1}) + b_{g2}\right) \quad (11)$$

Setting $\delta_{\text{floor}} = 0.10$ prunes 65.9% of spurious camouflage edge weights (Supplementary Fig. S3). During training on split $\mathcal{E}_{\text{train}}$, auxiliary supervision $\mathbf{g}_{ij}^{\text{target}} \in \{0, 1\}$ (Eq. 21) guides the gate using hub-degree heuristics ($\deg_{\text{out}}(j) \geq 50$ or rate $> 100 \text{ tx/day}$ receives 0.0, standard edges receive 1.0). At inference time, the trained MLP evaluates unseen edges strictly via multi-modal representations $\mathbf{c}_{ij}$ without external targets or static degree cutoffs.

Unified Spatio-Temporal Attention: Message passing integrates structural compatibility, temporal decay, and edge-trust

gating over degree-capped causal incoming neighborhoods ($\mathcal{N}_K^{\text{in}}(u), K \leq 15$):

$$\tilde{\alpha}_{uv}^{(t)} = \frac{\Omega(u, v, t)}{\sum_{w \in \mathcal{N}_K^{\text{in}}(u)} \Omega(u, w, t)}, \quad (12)$$

$$\Omega(u, v, t) = \exp\left(\frac{\mathbf{q}_u^T(\mathbf{k}_v + \mathbf{W}_E \mathbf{e}_{uv})}{\sqrt{d_k}}\right) \cdot w(\Delta t_{uv}) \cdot \hat{g}_{uv} \quad (13)$$

E. Typology-Clustered Latent GraphSMOTE & Flow Regularity

Design Rationale: Real-world financial crime encompasses distinct operational archetypes (smurfing dispersal, pooling fan-in, wash-trading loops, and pass-through layering). Naive Euclidean oversampling across the global minority class synthesizes hybrid instances that violate domain flow-balance regularity. C-STGB partitions minority illicit nodes strictly within the training split into latent typology clusters:

$$C_k = \{v \in \mathcal{V}_{\text{illicit}}^{\text{train}} \mid \cos(\mathbf{h}_v, \boldsymbol{\mu}_k) \geq \tau_{\text{sim}}\} \quad (14)$$

where centroids $\boldsymbol{\mu}_k$ ($K = 4$) are optimized via latent silhouette validation on $\mathcal{D}_{\text{train}}$. Synthetic representations are interpolated strictly along trajectories within the same cluster: $\mathbf{h}_{\text{syn}} = (1 - \rho)\mathbf{h}_u + \rho\mathbf{h}_v$, with $\rho \sim \mathcal{U}(0, 1)$ and $u, v \in C_k \subseteq \mathcal{V}_{\text{train}}$. Bilinear links to candidate training nodes $w \in \mathcal{V}_{\text{train}}$ are generated under topological edge reconstruction supervision:

$$\hat{\mathbf{A}}_{\text{syn}, w} = \sigma\left(\mathbf{h}_{\text{syn}}^T \mathbf{W}_{\text{edge}} \mathbf{h}_w\right) \cdot \mathbb{I}(t_w \leq t_{\text{syn}}) > \tau_{\text{edge}} \quad (15)$$

where $t_{\text{syn}} = \max(t_u, t_v)$ defines the causal birth timestamp, strictly precluding acausal edge formation with future transactions.

Preservation of Flow Conservation Regularity: Incoming links connect from candidate sources $w \in \mathcal{N}_{\text{in}}(u) \cup \mathcal{N}_{\text{in}}(v)$ and outgoing links connect to candidate destinations $w' \in \mathcal{N}_{\text{out}}(u) \cup \mathcal{N}_{\text{out}}(v)$. Synthetic amounts are parameterized linearly: $A_{\text{syn}} = (1 - \rho)A_u + \rho A_v$ and $B_{\text{syn}} = (1 - \rho)B_u + \rho B_v$. Under strictly positive fund flows ($A_u, A_v, B_u, B_v > 0$), the synthetic flow regularity ratio simplifies to a strict convex combination of parent ratios:

$$\Phi_{\text{flow}}(\mathbf{h}_{\text{syn}}) = \lambda \Phi_{\text{flow}}(u) + (1 - \lambda) \Phi_{\text{flow}}(v) \quad (16)$$

where $\lambda = \frac{(1 - \rho)B_u}{(1 - \rho)B_u + \rho B_v} \in (0, 1)$. Consequently, $\Phi_{\text{flow}}(\mathbf{h}_{\text{syn}})$ is guaranteed to lie within $[\min(\Phi_{\text{flow}}(u), \Phi_{\text{flow}}(v)), \max(\Phi_{\text{flow}}(u), \Phi_{\text{flow}}(v))]$, preserving domain mass-balance regularity without artificial distortion (formal proof in Supplementary Lemma S2).

F. Evidence-Adaptive Graph-Tabular Fusion (Cold-Start Safety Valve)

Design Rationale: In production payment rails, 38.4% of newly activated money mules possess sparse connectivity ($\deg(u) \leq 2$). Pure spatial message passing starves on isolated stubs (18.4% F1). C-STGB resolves this via an evidence-adaptive gating parameter $\alpha_u \in [0, 1]$ acting as a continuous topological confidence estimator:

$$\mathbf{h}_u^* = \alpha_u \mathbf{h}_u^{(L)} + (1 - \alpha_u) [\mathbf{z}_{\text{inv}}(u, t) \parallel \mathbf{x}_u], \quad (17)$$

$$\alpha_u = \sigma\left(\mathbf{w}_\alpha^T [\mathbf{h}_u^{(L)} \parallel \mathbf{z}_{\text{inv}}(u, t) \parallel \mathbf{x}_u] + b_\alpha\right) \quad (18)$$

When neighborhood evidence is sparse ($\deg(u) \leq 1$), $\alpha_u \rightarrow 0$ (empirically $\alpha_u \in [0.08, 0.12]$), gracefully transferring predictive mass to the 12-D canonical invariants $\mathbf{z}_{\text{inv}}(u, t)$ and achieving 89.60% F1 on cold-start stubs. Conversely, when local topology is dense, $\alpha_u \rightarrow 1.0$, harnessing high-order multi-hop structural representations (Supplementary Fig. S6).

G. Cost-Sensitive Optimization & Bayes Risk Minimization

Anomaly posterior probabilities evaluate as $\hat{p}_u = \sigma(\text{MLP}(\mathbf{h}_u^*))$. The multi-objective loss function is formulated as:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{task}}(y, \hat{p}) + \gamma_1 \mathcal{L}_{\text{recon}}(\mathbf{A}_{\text{train}}, \hat{\mathbf{A}}_{\text{train}}) + \gamma_2 \mathcal{L}_{\text{aux}} \quad (19)$$

where $\mathcal{L}_{\text{task}}$ is the Cost-Sensitive Asymmetric Focal Tversky Loss ($\alpha = 0.70, \beta = 0.30, \gamma = 1.5$), prioritizing minority recall while heavily penalizing false negatives. Bilinear edge reconstruction loss $\mathcal{L}_{\text{recon}}$ evaluates positive training edges $\mathcal{E}_{\text{train}}$ against a 1 : 5 negative non-edge sample $\mathcal{E}_{\text{neg}}$:

$$\begin{aligned} \mathcal{L}_{\text{recon}} = & -\frac{1}{|\mathcal{E}_{\text{train}}|} \sum_{(u, v) \in \mathcal{E}_{\text{train}}} \log \hat{A}_{uv} \\ & -\frac{1}{|\mathcal{E}_{\text{neg}}|} \sum_{(u, w) \in \mathcal{E}_{\text{neg}}} \log(1 - \hat{A}_{uw}) \end{aligned} \quad (20)$$

with $\hat{A}_{uv} = \sigma(\mathbf{h}_u^T \mathbf{W}_{\text{edge}} \mathbf{h}_v)$. Auxiliary loss $\mathcal{L}_{\text{aux}}$ enforces edge-trust sparsity and tri-band orthogonality ($\gamma_1 = 0.50, \gamma_2 = 0.10, \lambda_{\text{ortho}} = 0.05$):

$$\mathcal{L}_{\text{aux}} = \frac{1}{|\mathcal{E}_{\text{train}}|} \sum_{(i, j) \in \mathcal{E}_{\text{train}}} \|\mathbf{g}_{ij} - \mathbf{g}_{ij}^{\text{target}}\|_2^2 + \lambda_{\text{ortho}} \|\mathbf{W}_{\text{burst}}^T \mathbf{W}_{\text{seas}}\|_F^2 \quad (21)$$

Decision threshold τ^* is calibrated exclusively over validation split $\mathcal{D}_{\text{val}}$ minimizing operational Bayes risk under asymmetric penalty ratio $C_{\text{FN}}/C_{\text{FP}} = 15.0$:

$$\tau^* = \arg \min_{\tau \in [0, 1]} \mathbb{E}[C_{\text{FN}} \cdot \mathbb{I}(\hat{p} < \tau, Y = 1) + C_{\text{FP}} \cdot \mathbb{I}(\hat{p} \geq \tau, Y = 0)] \quad (22)$$

reducing operational risk by 78.4% relative to standard default thresholding ($\tau = 0.50$).

H. Class-Conditional Conformal Risk Control & 3-Tier Triage

Over independent holdout calibration split $\mathcal{D}_{\text{cal}}$, empirical class-conditional non-conformity quantiles $\hat{q}^{(y)}$ are computed ($y \in \{0, 1\}$):

$$\hat{q}^{(y)} = \text{Quantile}\left(s_1^{(y)}, \dots, s_{n(y)}^{(y)}; \left\lceil \frac{(n(y) + 1)(1 - \alpha_y)}{n(y)} \right\rceil\right) \quad (23)$$

where non-conformity scores are $s_i^{(y)} = 1 - \hat{p}_i(y)$.

Decoupled Theoretical Guarantees:

- 1) **Guarantee A (Exact Finite-Sample Validity):** Under within-class exchangeability on holdout split $\mathcal{D}_{\text{cal}}$, split conformal prediction guarantees exact class-conditional containment: $\mathbb{P}(Y \in \Gamma(X) \mid Y = y) \geq 1 - \alpha_y, \forall y \in \{0, 1\}$.
- 2) **Guarantee B (Non-Asymptotic Bound under Total Variation Drift):** Under streaming distribution shift between calibration $\mathcal{P}_{\text{cal}}^{(y)}$ and test $\mathcal{P}_t^{(y)}$, coverage degradation de-

Algorithm 1 C-STGB Unified Multi-Stage Pipeline

Require: Transaction graph $\mathcal{G}_t = (\mathcal{V}_t, \mathcal{E}_t)$, Error levels α_0, α_1 , Cost ratio $\frac{C_{FN}}{C_{FP}}$.

Ensure: Model parameters Θ , Threshold τ^* , Prediction set $\Gamma(u)$.

- 1: **Phase 1: Feature & Temporal Invariant Extraction**
- 2: Compute 12-D flow invariants $\mathbf{z}_{inv}(u, t)$ for $u \in \mathcal{V}_t$ via truncated power iteration taint and Hawkes arrival intensity $\lambda_u(t)$.
- 3: Compute continuous tri-band temporal embeddings $\Phi_{time}(\Delta t)$.
- 4: **Phase 2: Topology Filtering & Typology Oversampling**
- 5: Evaluate edge-trust gates $\hat{g}_{uv}$ (Eq. 10) and prune edges with $\hat{g}_{uv} < \delta_{floor}$.
- 6: Partition minority training nodes into $K = 4$ typology clusters C_k .
- 7: Synthesize latent representations $\mathbf{h}_{syn}$ and bilinear causal edges $\hat{\mathbf{A}}_{syn}$ preserving flow regularity within parent bounds.
- 8: **Phase 3: Deep Representation & Adaptive Fusion**
- 9: Aggregate multi-scale temporal messages over causal incoming neighbors $\mathcal{N}_K^in(u)$, $K \leq 15$ to obtain $\mathbf{h}_u^{(L)}$.
- 10: Compute evidence-adaptive gating α_u and fused state $\mathbf{h}_u^*$ (Eq. 17).
- 11: **Phase 4: Optimization & Conformal Calibration**
- 12: Optimize Θ minimizing multi-objective loss $\mathcal{L}_{total}$ (Eq. 19).
- 13: Calibrate decision threshold τ^* on validation split $\mathcal{D}_{val}$ (Eq. 22).
- 14: Compute class-conditional quantiles $\hat{q}^{(0)}, \hat{q}^{(1)}$ on calibration split $\mathcal{D}_{cal}$ (Eq. 23).
- 15: **Phase 5: Streaming Inference & Operational Triage**
- 16: Predict posterior risk $\hat{p}_u = \sigma(\text{MLP}(\mathbf{h}_u^*))$.
- 17: Construct conformal prediction set $\Gamma(u) = \{y \in \{0, 1\} : 1 - \hat{p}_u(y) \leq \hat{q}^{(y)}\}$.
- 18: **Route:** $\Gamma(u) = \{1\} \rightarrow \text{Tier 1}; \Gamma(u) = \{0, 1\} \rightarrow \text{Tier 2}; \Gamma(u) = \{0\} \rightarrow \text{Tier 3}$.

composes via maximal coupling into TV shift and empirical process concentration:

$$\begin{aligned} & |\mathbb{P}_t(Y \in \Gamma(X) \mid Y = y) - (1 - \alpha_y)| \\ & \leq d_{TV}(\mathcal{P}_t^{(y)}, \mathcal{P}_{cal}^{(y)}) + \sqrt{\frac{\ln(2/\delta)}{2n_y}} + \frac{1}{n_y} \end{aligned} \quad (24)$$

with probability $\geq 1 - \delta$, where $n_y = |\mathcal{D}_{cal}^{(y)}|$ (Proposition S1).

- 3) **Guarantee C (Long-Run Martingale Convergence under Delay):** Ground-truth confirmation incurs investigative delay $\Delta T_{delay} \in [30, 90]$ days. Streaming Adaptive Conformal Inference (ACI) updates online target quantiles over the confirmed observation queue:

$$\alpha_t \leftarrow \alpha_{t-1} + \gamma(\alpha - \text{err}_{t-\tau(t)}) \quad (25)$$

where $\text{err}_{t-\tau(t)} = \mathbb{I}(y_{t-\tau(t)} \notin \Gamma_{t-\tau(t)}(x_{t-\tau(t)}))$. By Martingale concentration for delayed feedback sequences [41], long-run marginal coverage converges almost surely: $\lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T \text{err}_t = \alpha$.

Operational 3-Tier Queue Mapping: Conformal prediction sets $\Gamma(X) \subseteq \{0, 1\}$ partition transactions into three operational workflows (Supplementary Fig. S1): (i) **Tier 1 (High-Risk Escalation):** $\Gamma(X) = \{1\}$, flagged for immediate compliance hold and automated SAR compilation; (ii) **Tier 2 (Selective Abstention Queue):** $\Gamma(X) = \{0, 1\}$, routed to compliance officers as ambiguous cases ($\approx 0.50\%$ of volume); and (iii) **Tier 3 (Straight-Through Low-Risk Screening Tier):** $\Gamma(X) = \{0\}$, cleared without human intervention ($> 99.4\%$ of volume).

I. Unified Algorithmic Specification & Systems Complexity

The complete end-to-end execution of C-STGB is formalized in Algorithm 1.

Computational Complexity and Streaming SLAs: Ingesting a streaming transaction requires: (1) DuckDB/Arrow

TABLE I
STRUCTURAL OVERVIEW OF THE 14 BENCHMARK FINANCIAL NETWORKS
(COMPLETE 14-DATASET ITEMIZATION IN SUPPLEMENTARY TABLE S3).

Archetype Group	Datasets Included	Entities	Txs	Illicit (π)
Group A (Crypto, 4)	Elliptic-v1/2, ETH, MtGox	2,621,048	3,430,755	0.09%–2.23%
Group B (Banking, 9)	SAML-D, PaySim, AMLSim, Gen, DGraph	6,629,125	5,812,793	0.05%–1.50%
Group C (Card, 1)	Credit Card Forensics	284,807	284,807	0.17%
Total Benchmark Corpus	14 Financial Networks	9,534,980	9,528,355	0.05%–2.23%

descriptor extraction (0.11 ms); (2) Top-15 causal subgraph retrieval (0.18 ms); (3) Temporal forward pass (0.12 ms); and (4) Conformal set routing (0.04 ms). The total streaming Fast-Path latency is **0.45** ms per event (P99: 0.82 ms) with 428 MB peak VRAM, sustaining **31,200** tx/s line-rate throughput and strictly bounded $O(K \cdot T_{iter})$ per-node complexity.

IV. EXPERIMENTAL EVALUATION

A. Benchmark Dataset Corpus & Topological Regimes

To evaluate model generalization across diverse financial rails, we benchmark C-STGB across 14 financial transaction networks spanning public cryptocurrency ledgers, multi-bank wire systems, and mobile financial services, summarized in Table I. The corpus comprises 9,534,980 entities and 9,528,355 transactions across three distinct topological archetypes:

- 1) **Group A (Public Blockchain Ledgers, 4 Networks):** Directed Acyclic Graphs (DAGs) and smart-contract transaction meshes (Elliptic-v1 [7], Elliptic-v2 [10], XBlock-ETH [16], MtGox [17]) characterized by global ledger observability and structural transparency.
- 2) **Group B (Multi-Bank, Mobile Synthetic & FinTech Rails, 9 Networks):** Inter-bank wire meshes, high-velocity mobile draining trees, and dynamic credit networks (SAML-D [18], PaySim1, PaySim Extended [45], IBM-AMLSim HI/LI Small and Medium (4 datasets) [19], the Synthetic Typology Data Generator, and DGraphFin [46]) characterized by extreme class imbalance ($\pi \in [0.05\%, 1.50\%]$) and fragmented organizational visibility. In `paysim_extended`, each mobile financial log record is mapped to a dedicated transaction-event vertex connected to its originating account and step context, yielding an exact 1:1 entity-to-edge projection ($|\mathcal{V}| = |\mathcal{E}| = 1,048,575$).
- 3) **Group C (Classical Tabular & Card Fraud Streams, 1 Network):** High-throughput card-terminal bipartite transaction streams (`cc_transactions` [47]).

B. Temporal Partitioning & Zero-Leakage Audit

To prevent temporal lookahead bias and evaluate realistic forward inductive deployment, all datasets follow a strict 4-way chronological partition: (1) **Parameter Training Split ($\mathcal{D}_{train}$, 60%):** used strictly for neural parameter updates; (2) **Validation Split ($\mathcal{D}_{val}$, 10%):** used for early stopping and Bayes threshold calibration (τ^*); (3) **Calibration Split ($\mathcal{D}_{cal}$, 10%):** dedicated holdout split for computing non-conformity quantiles $\hat{q}^{(0)}, \hat{q}^{(1)}$ (zero weight updates, zero threshold tuning); and (4) **Forward Out-of-Time Test Split ($\mathcal{D}_{test}$, 20%):** strictly future horizon for final evaluation.

For discrete UTXO graphs (e.g., Elliptic-v1 with 49 timesteps), Timesteps 1–30 (61.2%) form $\mathcal{D}_{train}$, Timesteps

31–34 (8.2%) form $\mathcal{D}_{\text{val}}$, Timesteps 35–38 (8.2%) form $\mathcal{D}_{\text{cal}}$, and Timesteps 39–49 (22.4%) serve as $\mathcal{D}_{\text{test}}$. Continuous transaction stream datasets (PaySim, SAML-D, IBM-AMLSim, Ethereum) follow an identical non-overlapping timestamp horizon partition ($t_0 \rightarrow t_{\text{train}} \rightarrow t_{\text{val}} \rightarrow t_{\text{cal}} \rightarrow t_{\text{test}}$).

Randomness Accounting across Seeds: While temporal split boundaries are strictly chronological and deterministic (60/10/10/20 partition), reported metrics reflect empirical variation across five locked random seeds ($S \in \{42, 101, 2024, 7, 999\}$). Seed variation governs four explicit stochastic components: (1) parameter initialization; (2) mini-batch shuffling trajectories; (3) latent GraphSMOTE stochastic interpolation coefficients $\rho \sim \mathcal{U}(0, 1)$; and (4) Monte Carlo dropout masks ($p = 0.20$).

Deterministic Topological Leakage Verification: A formal verification suite executes three topological invariant checks prior to training: (1) *Entity Boundary Isolation*: verifies $\mathcal{V}_{\text{test}} \cap \mathcal{V}_{\text{train}} = \emptyset$ for forward unseen accounts; (2) *Taint Masking Invariant*: enforces analytical taint seeds $\mathbf{s}_{\text{seed}}(u) \equiv 0$ for all $u \in \mathcal{D}_{\text{val}} \cup \mathcal{D}_{\text{cal}} \cup \mathcal{D}_{\text{test}}$; and (3) *Directional Time Invariance*: confirms that every edge $(u, v, t) \in \mathcal{E}_{\text{train}}$ strictly satisfies $t \leq t_{\text{train}}$. The verification suite confirms 0.00% entity overlap and zero future-timestamp propagation across all 14 empirical benchmarks. Precision-Recall and ROC curves on Elliptic-v1 across five locked seeds are provided in Supplementary Fig. S2.

C. Baseline Architectures & Fair Optimization Protocol

We benchmark C-STGB against 12 distinct baseline architectures across six canonical model paradigms: (1) *Tabular Tree Ensembles*: XGBoost [48], CatBoost [49], and LightGBM [50]; (2) *Spatial GNNs*: GCN [5], GraphSAGE [6], GAT [51], and GIN [52]; (3) *Heterogeneous GNNs*: Vanilla HGT [3] and Johannessen & Jullum [11]; (4) *Adversarial Fraud Detectors*: CARE-GNN [9]; (5) *Dynamic Snapshot Models*: EvolveGCN [8]; and (6) *Continuous Temporal GNNs*: TGN [22]. Evaluating these 12 baselines alongside C-STGB across 14 benchmark datasets yields 13 model configurations and exactly 182 dataset $\times$ model evaluation pairs, evaluated over five locked random seeds (910 total experimental executions).

Positioning Against Recent Flagship Literature: We contextualize C-STGB against three recent flagship works: (1) *H²IDE* (Fu *et al.*, *IEEE TKDE* 2025 [13]) isolates camouflaged fraudsters over multi-relational graphs via disentangled homophily and heterophily, but operates over static relations and lacks continuous-time streaming edge dynamics (Δt); (2) *CamFD* (Ou *et al.*, *IEEE TSMC* 2026 [29]) filters camouflaged fraudsters over discrete relational schemas without post-hoc finite-sample risk guarantees ($\geq 99\%$ coverage); and (3) *Group-Aware GNN* (Cheng *et al.*, *IEEE TKDE* 2023 [12]) clusters account groups on static ledgers rather than streaming transaction-level conformal calibration. As demonstrated in Section IV-J and Supplementary Table S2, C-STGB uniquely bridges continuous-time edge gating with distribution-free conformal triage.

Controlled Factorial Feature Attribution Protocol (Set A): To isolate whether empirical gains stem from the C-STGB

graph architecture, engineered domain features, or threshold optimization, we cross-evaluate models across five controlled feature and threshold regimes: (1) *Raw Features Only* ($\tau = 0.50$): GCN achieves $43.55 \pm 0.42\%$ F1 on `elliptic_v1` (14-dataset macro: $23.78 \pm 0.35\%$); (2) *Raw Features* (τ^*): GCN F1 improves to $51.04 \pm 0.38\%$ (+7.49 pp; macro: $28.12 \pm 0.40\%$), confirming threshold calibration is vital under extreme imbalance; (3) *Descriptors Only* ($\mathbf{z}_{\text{inv}}$): Tabular models on 12-D descriptors achieve $84.60 \pm 0.85\%$ F1 ($61.40 \pm 0.52\%$ macro), establishing an informative predictive floor; (4) *Raw Features + Descriptors* ($[\mathbf{x}_u \parallel \mathbf{z}_{\text{inv}}]$): Augmenting spatial GNNs yields marginal uplift: +1.27 pp on GCN (44.82%), +1.38 pp on GraphSAGE (50.45%), and +0.81 pp on EvolveGCN (51.85%); 14-dataset macro GCN reaches 24.85%, over 42 pp behind C-STGB’s 67.26%; full matrix in Supplementary Table S3); and (5) *Raw Features on Base C-STGB (No Invariants, $\tau = 0.50$)*: Base C-STGB achieves $79.08 \pm 0.55\%$ F1 on `elliptic_v1` ($48.20 \pm 0.50\%$ macro), substantially exceeding raw GCN (43.55%), GraphSAGE (49.07%), and TGN (64.10%).

Mechanistic Divergence Analysis: This controlled factorial comparison reveals a fundamental mechanistic divergence: isotropic spatial convolutions ($\tilde{\mathbf{D}}^{-1/2} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-1/2} \mathbf{X}$) linearly average feature representations across multi-hop neighborhoods. When illicit accounts connect to high-degree commercial hubs ($\text{deg} > 50$), localized scalar signatures in $\mathbf{z}_{\text{inv}}$ (e.g., flow regularity $\Phi_{\text{flow}} \approx 1.0$ and burst intensity λ_u) are arithmetically diluted by dozens of benign neighbors into the majority mean. Furthermore, lacking edge-trust gating ($\hat{g}_{ij}$) and latent GraphSMOTE, gradients under extreme imbalance ($\pi < 0.05\%$) are swamped by majority benign nodes. In contrast, tree ensembles evaluate axis-aligned splits on unmixed features, isolating Φ_{flow} directly. Thus, extracting relational multi-hop patterns without over-smoothing strictly requires C-STGB’s integrated architecture. All architectures operate under calibrated parameter budgets (1.84×10^6 for C-STGB vs. 1.80×10^6 for HGT and 2.12×10^6 for TGN; $N = 50$ Bayesian trials with identical tuning grids).

D. Evaluation Metrics & Statistical Testing Protocol

Under extreme class imbalance ($\pi \ll 0.01$), traditional ROC-AUC is overly optimistic due to the dominant true negative denominator, and binary macro F1 arithmetic means are artificially dominated by the $> 99\%$ benign class ($F_1^{(0)} \approx 1.0$). We evaluate: (1) **Precision-Recall AUC (PR-AUC)**: threshold-independent minority trade-offs (Supplementary Fig. S2); (2) **Minority (Illicit) Class F1-Score ($F_1^{(1)}$) & Recall**: harmonic utility on confirmed illicit entities under validation-tuned threshold τ^* ; and (3) **Conformal Set Efficiency** ($\mathbb{E}[|\Gamma(X)|]$): prediction set tightness. Across the corpus, Macro-Averages denote the unweighted arithmetic mean of minority F1 across all 14 datasets ($\frac{1}{14} \sum_{i=1}^{14} F_1^{(1)}(i)$). All models undergo paired threshold optimization on $\mathcal{D}_{\text{val}}$ under identical search grids minimizing $\mathcal{L}_{\text{ops}}$ ($C_{\text{FN}}/C_{\text{FP}} = 15.0$). Statistical significance is evaluated via two-sided Wilcoxon Signed-Rank Tests across all 14 benchmark networks, with Benjamini–Hochberg False Discovery Rate (FDR) multiple-testing correction.

E. Hardware Compute Infrastructure

All experiments were conducted across Kaggle Cloud TPU v3-8 / NVIDIA Tesla T4 GPUs (16GB VRAM) and a dedicated workstation comprising an AMD Ryzen 9 5900X CPU (12 cores, 24 threads), 64GB RAM, and an NVIDIA RTX 3080 GPU (10GB VRAM) running Ubuntu 22.04 LTS (14.2 GPU-hours local, 18.5 GPU-hours cloud).

F. Comparative Baseline Performance & Benchmark Scorecard

Table II reports benchmark comparisons across all 14 benchmark networks (with the complete multi-metric scorecard and per-baseline profiling detailed in Supplementary Table S3). Across the corpus, C-STGB achieves 99.44% F1 on Elliptic-v1 (PR-AUC 1.0000), 100.00% F1 on Elliptic-v2 (PR-AUC 1.0000), 99.80% on PaySim Extended, 97.91% on DGraphFin, and 96.94% on XBlock-ETH.

Statistical Parity with Trees vs. Superiority Over Deep GNNs: Across the 14-dataset benchmark suite, C-STGB yields an unweighted macro F1 of 67.26% (0.6848 PR-AUC). Direct comparison against tuned gradient-boosted decision trees reveals statistical parity on the global average: against XGBoost, C-STGB exhibits a -0.66 pp global difference (67.26% vs. 67.92%) with 6 wins, 1 tie, and 7 losses ($W = 48.0$, $p_{\text{adj}} = 0.8077$, n.s.); against CatBoost, C-STGB exhibits a $+0.13$ pp global difference (67.26% vs. 67.14%) with 7 wins, 2 ties, and 5 losses ($W = 32.0$, $p_{\text{adj}} = 0.4318$, n.s.). Therefore, C-STGB does not claim universal statistical superiority over strong tabular learners globally. Rather, empirical evidence demonstrates that C-STGB delivers statistically significant outperformance over all evaluated deep spatial and dynamic GNN baselines: $+43.49$ pp mean uplift over GCN (median $+37.11$ pp, IQR 72.66 pp; 14 wins, 0 losses), $+42.95$ pp over GraphSAGE (median $+37.07$ pp, IQR 69.08 pp; 14 wins, 0 losses), and $+48.32$ pp over EvolveGCN (median $+41.91$ pp, IQR 56.75 pp; 14 wins, 0 losses) ($W = 105.0$, $W_- = 0.0$, $n = 14$, $p = 1.22 \times 10^{-4}$, $p_{\text{adj}} < 0.001$). Group-macro averages reflect this domain divergence: Group A Crypto (92.15%), Group B Banking/FinTech (58.07%), and Group C Card (51.40%), with median dataset F1 of 61.81% (IQR 74.72 pp). The wide interquartile range highlights an intrinsic bimodal partition across the benchmark corpus: Cluster 1 comprises publicly indexed cryptocurrency ledgers with dense Darknet market seizures exhibiting high structural separability ($> 96\%$ F1 across modern paradigms), whereas Cluster 2 comprises complex multi-bank synthetic laundering simulations (e.g., PaySim1, IBM-AMLSim LI) where performance drops below 45%, representing the true open challenge in adversarial financial forensics.

When Graph Learning Helps — and When It Does Not: This performance pattern highlights a critical domain boundary: (1) *Multi-Hop Relational Regimes*: On structurally complex networks with multi-hop laundering chains, high degree variance, and camouflage injection (e.g., IBM-AMLSim HI, SAML-D, Elliptic-v1/v2), C-STGB outperforms both base GNNs ($+9.98$ pp to $+48.3$ pp) and tree ensembles ($+1.05$ pp to $+4.67$ pp on IBM-AMLSim HI), as spatio-temporal message

passing captures coordinated fund movement across account chains that cannot be represented in flat tabular rows. (2) *Flat Transaction Streams*: Conversely, on paysim1, C-STGB achieves only 10.68% F1, trailing XGBoost (18.90%) and CatBoost (17.76%). paysim1 is a synthetic mobile payment simulation where fraud signals reside in localized tabular balance discrepancies within single isolated transactions. Without multi-hop laundering chains, neighborhood graph aggregation provides no structural advantage over axis-aligned decision trees.

Forensic Investigation of Near-Perfect Benchmark Scores: Multiple model families achieve near-perfect scores on Elliptic-v1 (99.44% C-STGB, 99.89% XGBoost, 99.78% CatBoost) and Elliptic-v2 (100.00% C-STGB, 99.81% XGBoost, 100.00% CatBoost). Forensic examination reveals that these scores stem from the underlying structure of public Bitcoin UTXO benchmarks: ground-truth illicit entities originate from law enforcement Darknet market seizures (e.g., Silk Road, AlphaBay), which form dense, topologically segregated connected subgraphs with sharp boundaries and near-perfect label homophily ($h > 0.92$). A linear probe sanity check confirms high intrinsic separability: Topological Logistic Regression achieves 83.1% F1 on Elliptic-v2, verifying that these benchmarks exhibit high structural separability rather than presenting an unconstrained real-world laundering challenge.

Statistical Significance Testing: Table III summarizes non-parametric hypothesis testing with dataset-level paired observations across all $n = 14$ benchmarks: two-sided Wilcoxon signed-rank tests yield maximal positive rank sum $W = 105.0$ ($W_- = 0.0$, $n = 14$), with exact two-sided $p = 1.22 \times 10^{-4}$ and Benjamini-Hochberg adjusted $p_{\text{adj}} = 6.10 \times 10^{-4} < 0.001$ against GCN, GraphSAGE, and EvolveGCN, reflecting uniform outperformance across all 14 empirical networks. Re-testing against invariant-augmented GNNs ($[\mathbf{x}_u \parallel \mathbf{z}_{\text{inv}}]$) preserves identical significance ($W = 105.0$, 14 wins, 0 losses, $p_{\text{adj}} < 0.001$; full 14-dataset paired scores in Supplementary Table S3).

G. Class-Conditional Conformal Risk Control & Queue Triage

Table IV evaluates Class-Conditional Conformal Risk Control (CRC) across benchmark archetypes at significance level $\alpha = 0.01$ (99.0% target true label coverage, full evaluation across calibration baselines in Supplementary Table S4). Across all datasets, empirical coverage satisfies the nominal 99.0% bound (99.02% – 99.28%), while routing $> 99.4\%$ of transaction volume through straight-through Tier 1 (High-Risk Escalation) and Tier 3 (Straight-Through Low-Risk Screening Tier) channels.

Crucially, the high coverage guarantee (99.12% macro-average) is achieved with prediction set efficiency: across all 9.53 million transactions, singleton decisive prediction sets ($|\Gamma(X)| = 1$) account for 99.50% of volume (Tier 1 + Tier 3), while ambiguous abstention sets ($|\Gamma(X)| = 2$, Tier 2) comprise only 0.50%. The empty prediction set rate $\mathbb{P}(\Gamma(X) = \emptyset)$ is verified to be exactly 0.00%, yielding an average prediction set size of $\mathbb{E}[|\Gamma(X)|] = 1.0050$, confirming that coverage guarantees are operationally tight.




Fig. 2. Empirical representation and classification benchmarks: (a) Precision-Recall and ROC curves demonstrating discrimination under extreme class skew; (b) t-SNE projection of the learned latent manifold under Typology GraphSMOTE, showing dense clustering of confirmed illicit topologies ($K = 4$) with clear separation from majority commercial hubs.

TABLE II

COMPREHENSIVE BENCHMARK COMPARISON ACROSS ALL 14 FINANCIAL NETWORKS: MINORITY (ILLICIT) CLASS F1-SCORE (%) AND PR-AUC WITH VALIDATION-TUNED THRESHOLDS τ^* UNDER ASYMMETRIC OPERATIONAL LOSS AGAINST LITERATURE BASELINES (TABLE DISPLAYS 5 REPRESENTATIVE BASELINES ALONGSIDE C-STGB DUE TO LAYOUT WIDTH; THE FULL MULTI-BASELINE BENCHMARK MATRIX ACROSS TABULAR, SPATIAL, AND DYNAMIC PARADIGMS IS REPORTED IN SUPPLEMENTARY TABLE S3; $\dagger$ INDICATES STATISTICALLY SIGNIFICANT CORPUS-WIDE IMPROVEMENT OVER DEEP GNNs AT $W = 105.0$, $p_{\text{adj}} < 0.001$ UNDER TWO-SIDED WILCOXON SIGNED-RANK TEST ACROSS ALL 14 BENCHMARK NETWORKS).

Dataset Identifier	XGBoost (Tabular)	GCN (Spatial)	EvolveGCN (Dynamic)	GraphSAGE (Inductive)	CatBoost (Industrial)	C-STGB (Proposed)
Group A: Public Blockchain Ledgers & Relational DAGs (4 Networks)						
elliptic_v1	99.89/1.0000	43.55/0.3482	51.04/0.4243	49.07/0.4254	99.78/1.0000	99.44/1.0000
elliptic_v2	99.81/0.9973	0.00/0.0237	0.00/0.0258	0.24/0.0230	100.00/1.0000	100.00/1.0000
xblock_eth	96.91/0.9961	6.53/0.0212	6.62/0.0211	6.75/0.0206	95.91/0.9925	96.94/0.9915
mtgox_leaked	74.39/0.8229	20.30/0.2038	22.12/0.0956	22.16/0.1801	71.87/0.7983	72.21/0.8306
Group B: Multi-Bank, Mobile Synthetic & FinTech Rails (9 Networks)						
saml_d	93.74/0.9593	1.69/0.0109	3.81/0.0077	3.16/0.0071	93.60/0.9448	93.68/0.9574
paysim1	18.90/0.1254	3.50/0.0084	3.98/0.0097	0.56/0.0016	17.76/0.1061	10.68/0.1173
paysim_extended	99.72/0.9996	96.22/0.9825	92.64/0.7521	96.05/0.9799	99.77/0.9996	99.80/0.9993
ibm_amlsim_hi_small	36.66/0.3407	2.46/0.0101	2.31/0.0080	2.46/0.0080	33.03/0.3100	37.70/0.3550
ibm_amlsim_hi_medium	42.97/0.4513	3.88/0.0135	7.06/0.0350	3.95/0.0126	41.53/0.4220	42.85/0.4609
ibm_amlsim_li_small	15.31/0.1350	0.84/0.0060	1.18/0.0052	1.50/0.0057	13.54/0.0802	15.76/0.1485
ibm_amlsim_li_medium	23.19/0.1811	3.41/0.0143	4.29/0.0238	2.30/0.0073	22.40/0.1718	23.38/0.2077
data_generator	100.00/1.0000	99.74/0.9979	62.72/0.6327	99.63/0.9956	100.00/1.0000	99.93/1.0000
dgraphfin	98.23/0.9978	2.60/0.0129	2.59/0.0094	3.26/0.0161	98.45/0.9980	97.91/0.9979
Group C: Bipartite Card Transaction Streams (1 Network)						
cc_transactions	51.15/0.5126	48.16/0.3472	4.79/0.0215	49.24/0.3555	52.26/0.5092	51.40/0.5213
Group A Macro (Crypto, 4 Sets)						
	92.75/0.9541	17.60/0.1492	19.95/0.1417	19.56/0.1623	91.89/0.9477	92.15/0.9531
Group B Macro (Banking/FinTech, 9 Sets)						
	58.83/0.5774	23.47/0.2227	20.00/0.1633	23.65/0.2201	58.01/0.5617	58.07/0.5843
Group C Macro (Card Stream, 1 Set)						
	51.15/0.5126	48.16/0.3472	4.79/0.0215	49.24/0.3555	52.26/0.5092	51.40/0.5213
Overall Macro Median [IQR]	67.68 [74.65]	3.69 [37.28]	4.54 [39.88]	3.61 [38.40]	67.14 [74.72]	61.81 [74.72]
Overall Macro Mean (All 14 Sets)	67.92/0.6799	23.78/0.2143	18.94/0.1480	24.31/0.2170	67.14/0.6666	67.26/0.6848[†]

Minority F1-Score evaluates harmonic precision and recall on confirmed illicit transactions ($y = 1$) under validation-calibrated threshold τ^* ; Macro-Average represents the unweighted arithmetic mean across all 14 dataset-level scores ($\frac{1}{14} \sum_{i=1}^{14} \text{Metric}_i$), alongside median and interquartile range (IQR, in percentage points). Representative canonical baselines are shown; complete profiling for all 12 baselines (including LightGBM with macro F1 66.62%, PR-AUC 0.6917) is itemized in Supplementary Table S3. Conceptual and taxonomic positioning against recent flagship literature (including H²IDE [13] and CamFD [29]) is detailed in Section IV-C and Supplementary Table S2. $\dagger$ indicates statistically significant corpus-wide improvement over deep GNNs at $W = 105.0$, $p_{\text{adj}} < 0.001$ under two-sided Wilcoxon signed-rank test across all 14 benchmark networks.

Group A: 4 Public Crypto Ledgers; Group B: 9 Multi-Bank, Mobile & FinTech Networks; Group C: 1 Bipartite Card Stream.

TABLE III

NON-PARAMETRIC STATISTICAL SIGNIFICANCE AND DISTRIBUTIONAL EFFECT SIZES OF C-STGB VS. BASELINES ACROSS ALL $n = 14$ BENCHMARK NETWORKS (TWO-SIDED WILCOXON SIGNED-RANK TEST WITH DATASET-LEVEL PAIRED OBSERVATIONS AND BENJAMINI-HOCHBERG FDR CORRECTION).

Baseline	W / T / L	Mean Δ	Median [IQR]	W Stat	Adjusted p_{adj}
XGBoost (Tabular)	6 / 1 / 7	-0.66 pp	-0.02 [0.46] pp	48.0	0.8077 (n.s.)
CatBoost (Industrial)	7 / 2 / 5	+0.13 pp	+0.05 [1.28] pp	32.0	0.4318 (n.s.)
GCN (Spatial)	14 / 0 / 0	+43.49 pp	+37.11 [72.66] pp	105.0	6.10×10^{-4}
GraphSAGE (Inductive)	14 / 0 / 0	+42.95 pp	+37.07 [69.08] pp	105.0	6.10×10^{-4}
EvolveGCN (Dynamic)	14 / 0 / 0	+48.32 pp	+41.91 [56.75] pp	105.0	6.10×10^{-4}

*Under invariant representations $[\mathbf{x}_u \| \mathbf{z}_{\text{inv}}]$, C-STGB preserves 14/0/0 wins ($W = 105.0$, $p_{\text{adj}} < 0.001$; full 14-dataset paired matrix in Supplementary Table S3).

TABLE IV

CLASS-CONDITIONAL CONFORMAL RISK CONTROL: COVERAGE AND TRIAGE WORKLOAD AT $\alpha = 0.01$ (99.0% TARGET CONTAINMENT, EXTENDED DYNAMICS IN SUPPLEMENTARY TABLE S4).

Dataset	Empirical Cov. ($\geq 99.0\%$ Target)	Selective Risk (Tier 3 FN)	Review Rate (Tier 2 Abstain)	Efficiency (Workload Red.)
elliptic_v1	99.12%	0.08%	0.48% (98)	-99.52%
elliptic_v2	99.05%	0.80%	0.55% (67)	-99.45%
dgraphfin	99.18%	0.13%	0.42% (1,470)	-99.58%
paysim_extended	99.22%	0.12%	0.38% (797)	-99.62%
saml_d	99.15%	0.05%	0.45% (653)	-99.55%
cc_transactions	99.10%	0.15%	0.58% (330)	-99.42%
ibm_amlsim_hi_med	99.08%	0.22%	0.52% (286)	-99.48%
Macro (All 14 Sets)	99.12%	0.18%	0.50%	-99.50%

H. Modular Stepwise & Leave-One-Out Ablation Study

To evaluate architectural contributions, Table V details both cumulative stepwise progression from baseline static GNN to full C-STGB across three financial archetypes (Bitcoin UTXO, Mobile Money, Multi-Bank rails) and leave-one-out removals on Elliptic-v1. Baseline GCN achieves 43.55% F1 at $\tau = 0.50$ (Step 1a) and 51.04% at τ^* (Step 1b). Adding Continuous Temporal Encoding (+Mod 1, 68.50%), Edge Trust Gating (+Mod 2, 81.20%), Typology GraphSMOTE (+Mod 3, 93.60%), Adaptive Fusion (+Mod 4, 97.80%), and Cost-Sensitive Thresholding (+Mod 5) progressively drives performance to 99.44% F1 (Recall 98.88%, Precision 100.00%, PR-AUC 1.0000). Conformal Risk Control (+Mod 6) delivers uncertainty triage without altering point F1. Leave-one-out removals confirm Cost-Sensitive Thresholding (-12.28 pp) and Typology GraphSMOTE (-10.39 pp) provide the largest individual margins.

Cold-Start ($d \leq 2$) Dissection & Gating Safety-Valve Dynamics: A vital forensic query is whether graph neural message passing adds value on sparsely connected accounts ($d_u \leq 2$), which constitute 38.4% of newly activated money mules in the Elliptic-v1 benchmark. When evaluated

TABLE V
UNIFIED ARCHITECTURAL ABLATION STUDY: (A) CUMULATIVE STEPWISE FORWARD PROGRESSION; (B) LEAVE-ONE-OUT COMPONENT REMOVALS ON ELLIPTIC-V1 (5 SEEDS).

(a) Cumulative Stepwise Forward Progression (Recall / F1-Score (%) across 5 Seeds)						
Ablation Progression	Elliptic-v1		PaySim Extended		IBM-AMLSim HI	
	Rec.	F1	Rec.	F1	Rec.	F1
(1a) Base Static GNN ($\tau = 0.50$)	34.20	43.55	84.20	89.50	4.10	7.85
(1b) Base Static GNN (τ^*)	42.50	51.04	91.20	94.80	6.80	11.50
(2) + Mod 1: Continuous Time	64.10	68.50	94.50	96.20	11.20	16.30
(3) + Mod 2: Edge Trust Gate	78.40	81.20	96.80	97.50	14.80	21.40
(4) + Mod 3: Typology SMOTE	91.50	93.60	98.10	98.60	19.50	28.60
(5) + Mod 4: Adaptive Fusion	96.80	97.80	99.10	99.30	23.10	33.80
(6a) + Mod 5: Cost-Sensitive τ^*	98.88	99.44	99.65	99.80	25.55	37.70
(6b) + Mod 6: Conformal Risk Control	<i>Preserves point F1; delivers 99.12% coverage with 1.005 avg set size</i>					
(b) Leave-One-Out Component Removal on Elliptic-v1 (Margin vs. Full Architecture)						
Ablated Component	Recall (%)	Precision (%)	F1-Score (%)			
Full C-STGB Architecture	98.88 ± 0.15	100.00 ± 0.00	99.44 ± 0.10			
w/o Typology GraphSMOTE	84.20 ± 0.85	94.50 ± 0.60	89.05 ± 0.70 (−10.39 pp)			
w/o Edge-Gating Filter (g_{ij})	91.50 ± 0.70	95.10 ± 0.55	93.25 ± 0.60 (−6.19 pp)			
w/o Tri-Band Temporal Decay	93.10 ± 0.65	96.80 ± 0.50	94.90 ± 0.55 (−4.54 pp)			
w/o Evidence-Adaptive Fusion (α_u)	95.40 ± 0.50	98.10 ± 0.40	96.72 ± 0.45 (−2.72 pp)			
w/o Hawkes Point Process Intensity	95.80 ± 0.45	98.40 ± 0.35	97.08 ± 0.40 (−2.36 pp)			
w/o 12-D Domain Descriptors (z_{dom})	96.80 ± 0.45	98.90 ± 0.35	97.84 ± 0.40 (−1.60 pp)			
w/o Cost-Sensitive Threshold (τ^*)	82.40 ± 0.90	92.50 ± 0.75	87.16 ± 0.80 (−12.28 pp)			

strictly on this cold-start subset on Elliptic-v1: (1) Standalone Hawkes-HGT suffers severe structural starvation, achieving only $18.40 \pm 1.15\%$ F1 due to degenerate neighborhood message passing; (2) Tabular GBDT over 12-D domain descriptors alone yields $84.60 \pm 0.85\%$ F1, capitalizing on flow balance and burst acceleration; and (3) C-STGB achieves **$89.60 \pm 0.62\%$** F1. Inspection of the learned gating parameter reveals that the gating network assigns $\alpha_u \in [0.08, 0.12]$ on cold-start stubs, smoothly transferring predictive mass to $\phi(\mathbf{z}_{\text{inv}}(u))$ while extracting subtle temporal signals from 1-hop seed transactions (t-SNE latent manifold separation is illustrated in Fig. 2b). Thus, the degree gate functions as a continuous topological safety valve preventing the catastrophic degradation typical of pure graph architectures on sparse topologies.

I. Latency-Memory Pareto Frontier & Systems Disentanglement

To avoid confounding model capacity with operational latency, we explicitly distinguish two operational configurations:

- 1) C-STGB-Fast (Streaming Triage Configuration):** Designed for wire-speed inline transaction scoring. Evaluates degree-capped 1-hop causal neighborhoods ($K \leq 15$) with pre-computed cached DuckDB descriptors. Achieves a streaming single-event latency of 0.45 ms (P99: 0.82 ms) and 428 MB peak VRAM (31, 200 tx/s, Table VI), itemized into: (i) DuckDB/Arrow descriptor lookup (0.11 ms); (ii) Top-15 causal subgraph extraction (0.18 ms); (iii) Temporal GNN forward pass (0.12 ms; 0.054 ms amortized under batch-64); and (iv) Conformal triage routing (0.04 ms, Supplementary Table S6).
- 2) C-STGB-Full (Full Forensic Configuration):** Employs the complete multi-scale tri-band attention with 2-hop neighborhood expansion, generating the benchmark accuracy results reported in Table II. Under this configuration, single-event streaming latency is 2.10 ms (3.30 ms batch-64) with 894 MB peak VRAM and 19,400 tx/s throughput.

J. Adversarial Threat Model & Topological Camouflage Defense

Formal Threat Model Specification: We formalize the adversarial evasion setting via four attributes:

TABLE VI
OPERATIONAL SYSTEMS SCALING AND LATENCY PROFILE ACROSS MODEL CONFIGURATIONS ON BENCHMARK HARDWARE.

Configuration	Streaming (P99)	Batch-64	Throughput	Peak VRAM
XGBoost (Tabular)	0.18 ms	0.95 ms	68,400 tx/s	112 MB
C-STGB (Fast-Path)	0.45 ms (0.82)	2.10 ms	31,200 tx/s	428 MB
C-STGB (Full Tri-Band)	2.10 ms (3.30)	3.30 ms	19,400 tx/s	894 MB
Vanilla HGT (2020)	8.40 ms	18.20 ms	3,500 tx/s	1,420 MB
TGN (Rossi <i>et al.</i>)	22.40 ms	41.80 ms	1,520 tx/s	3,280 MB

- 1) Attacker Knowledge:** Black-box (topology manipulation without model parameters) vs. White-box (full access to model parameters, gradients, and edge gating weights).
- 2) Attacker Budget:** Perturbation ratio bounded by $\|\delta\|_0 \leq 0.50|\mathcal{E}_u|$ for edge additions, amount perturbation $\|\Delta A\| \leq 0.20A$, and temporal shift $\Delta t \leq 72$ hours.
- 3) Feasible Attacker Actions:** (i) Random noise injection; (ii) Degree-matched hub camouflage; (iii) Low-degree camouflage (connecting to low-degree retail accounts); (iv) Temporal dilation; and (v) White-box projected gradient descent (PGD) topology rewiring.
- 4) Attacker Objective:** Minimize predicted risk score $\hat{p}_u < \tau^*$, evading detection while completing fund transfers.

Empirical Robustness Results & Construct Overlap Disclosure: Fig. 3 evaluates robustness across progressive perturbation ratios $\rho \in [0, 0.80]$ with 95% confidence bands across five seeds. Standard GCN collapses to $6.10 \pm 0.85\%$ F1 under gradient attacks and Vanilla HGT degrades to $12.40 \pm 1.10\%$ F1, whereas C-STGB retains $93.40 \pm 0.65\%$ F1.

We explicitly disclose an inductive construct overlap between Module 2 training supervision and degree-matched camouflage testing: auxiliary edge-trust loss (Eq. 21) guides $\hat{g}_{ij}$ during training using hub-degree heuristics ($\text{deg}_{\text{out}}(j) \geq 50$) on $\mathcal{E}_{\text{train}}$, aligning with the structural pattern manipulated in degree-matched camouflage. Consequently, retention under degree camouflage partly reflects this intended architectural bias.

To rigorously validate defense capabilities beyond this heuristic, we evaluate attacks completely outside the degree-matched construct: (1) *White-box gradient topology attacks:* where perturbations are guided strictly by $\nabla_{\mathcal{E}} \mathcal{L}$ without targeting node degrees, where C-STGB retains 93.40% F1; (2) *Low-degree camouflage:* connecting to low-degree accounts, where C-STGB retains 94.10% F1 due to flow-balance descriptors; and (3) *Temporal delay camouflage:* diluting velocity over weeks, where tri-band harmonic attention preserves 92.80% F1. Edge-trust pruning achieves 91.4% precision, 88.7% recall (AUROC 0.942), and false-prunes only 2.1% of legitimate commercial edges.

K. Hyperparameter Sensitivity & Stability Analysis

We evaluate operational stability across four key hyperparameters (Supplementary Table S5 and Supplementary Table S7): (1) *Edge-Trust Threshold* ($\delta_{\text{floor}} \in [0.02, 0.30]$): Macro F1 remains stable across $\delta_{\text{floor}} \in [0.06, 0.16]$ (± 0.35 pp around 99.44%). (2) *Operational Cost Ratio* ($C_{\text{FN}}/C_{\text{FP}} \in [5, 30]$): Calibrated decision thresholds adapt monotonically from $\tau^* = 0.42$ ($C_{\text{FN}}/C_{\text{FP}} = 5$) to $\tau^* = 0.18$ ($C_{\text{FN}}/C_{\text{FP}} = 30$).




Fig. 3. Adversarial Camouflage Robustness: Macro F1 retention across perturbation ratios $\rho \in [0, 0.80]$ on Elliptic-v1 under degree-matched and gradient-based evasion attacks (95% confidence bands, 5 seeds).

Expected operational loss remains within 4.2% of optimum across $[10, 20]$. (3) *Degree Cap* ($K \in [5, 30]$): F1 plateaus at $K = 15$ (99.44%), capturing 98.6% of multi-hop laundering while bounding retrieval latency to 0.18 ms (2.14 ms at $K = 50$). (4) *Adaptive Conformal Step Size* ($\gamma \in [0.001, 0.05]$): Coverage remains bounded within $[98.83\%, 99.34\%]$, with $\gamma = 0.005$ achieving optimal tracking stability.

V. DISCUSSION & OPERATIONAL IMPLICATIONS

A. Architectural Synergy & Division of Labor Across Modules:

Empirical evaluations across 14 benchmark networks yield five operational insights. Ablations (Table V) confirm performance stems from structured module specialization: (1) **Continuous Temporal Attention** captures microsecond bursts alongside 90-day dormancy (43.55% $\rightarrow$ 68.50% F1 on Elliptic-v1); (2) **Camouflage Filtration** prunes 65.9% of injected camouflage links ($g_{ij} < 0.10$), preserving a +5.62 pp margin; (3) **Latent Typology GraphSMOTE** elevates GNN recall from 49.07% to 98.88% (+5.64 pp F1 to reach 99.44% on Elliptic-v1 under invariant Bayes thresholding, vs. 91.42% graph-only in Supplementary Table S5); and (4) **Evidence-Adaptive Gating** acts as a safety valve: on dense graphs ($d \geq 5$), the GNN captures multi-hop loops, while on sparse stubs ($d \leq 2$), weight shifts to tabular descriptors ($\mathbf{z}_{inv}$), securing 89.60% F1 on cold-start entities (Supplementary Table S3).

B. Relational Topology vs. Tabular Tree Baseline Nuance:

Across 14 benchmarks, C-STGB achieves 67.26% unweighted macro F1 (0.6848 PR-AUC), demonstrating decisive superiority over evaluated deep GNNs (GCN 23.78%, GraphSAGE 24.31%, EvolveGCN 18.94%, all $p_{adj} < 0.001$). Comparison against tuned gradient-boosted decision trees reveals parity on global macro-average: C-STGB achieves 67.26% vs. CatBoost 67.14% ($p_{adj} = 0.4318$) and XGBoost 67.92% ($p_{adj} = 0.8077$). This demarcates a domain boundary: tree ensembles excel on single-hop tabular data (e.g., PaySim1, XGBoost 18.90% vs. C-STGB 10.68%) where fraud signals reside in balance discrepancies without relational graph structure. Conversely,

on multi-hop laundering topologies, C-STGB provides decisive advantages (+9.98 pp over base GNNs and +1.05 pp to +4.67 pp over trees on IBM-AMLSim HI), confirming spatio-temporal message passing is indispensable for coordinated transaction flows. In high-throughput industrial payment rails, this motivates a hierarchical production pattern: screening trivial single-hop transactions via low-cost trees, while escalating accounts exhibiting multi-hop burstiness, cold-start stubs, or suspected camouflage to C-STGB.

C. Cross-Domain Visibility (Crypto Ledgers vs. Banking Silos): C-STGB reaches 99.44% F1 on Elliptic-v1, 100.00% on Elliptic-v2, 96.94% on XBlock-ETH, and 99.80% on PaySim Extended because public blockchains and centralized ledgers provide globally observable transaction graphs. Conversely, on fragmented multi-bank rails (IBM-AMLSim, SAML-D), performance settles within 15.76%–42.85% F1 on IBM-AMLSim and 93.68% on SAML-D because individual institutions observe only local 1-hop transfers, blinding single-bank models to multi-bank layering and motivating federated graph architectures.

D. Economic Impact of Asymmetric Bayes Risk Policy:

In AML monitoring, undetected false negatives expose institutions to severe statutory penalties, whereas false positives consume officer review resources ($\approx$ \$50–\$150 per alert). To model this asymmetry, we adopt an operational loss proxy $\mathcal{L}_{ops} = C_{FN} \cdot FN + C_{FP} \cdot FP$ with penalty ratio $C_{FN}/C_{FP} = 15.0$. Optimizing Bayes decision threshold τ^* under Asymmetric Focal Tversky Loss reduces expected operational loss by **78.4%** relative to default thresholding ($\tau = 0.50$, Supplementary Table S4). The ratio $C_{FN}/C_{FP} = 15.0$ serves as an operational risk proxy that institutions can dynamically re-tune as enforcement priorities evolve.

E. Conformal Risk Control as an Operational Governance Paradigm:

Under model governance frameworks (including Federal Reserve SR 26-2 [53] and EU AI Act auditability mandates), institutions must bound decision uncertainty under operational shift. C-STGB provides finite-sample class-conditional coverage ($\geq 99.0\%$) under within-class exchangeability (Guarantee A), augmented by drift bounds (Guarantee B) and streaming Adaptive Conformal Inference (Guarantee C). Conformal coverage bounds the probability of true label containment under stated mathematical assumptions; it does not confer legal absolution. By isolating ambiguous transactions ($\Gamma(X) = \{0, 1\}$) into Tier 2 for compliance review while maintaining average set size $\mathbb{E}[|\Gamma(X)|] = 1.0050$ with 0.0% empty sets, C-STGB routes $> 99.4\%$ of transactions through a straight-through screening tier, harmonizing throughput with statistically grounded governance.

VI. QUALITATIVE FORENSIC CASE STUDY & ASSISTED COMPLIANCE DECISION SUPPORT

A. 15-Node Fan-In & Peeling Chain Model Attribution:

To evaluate model interpretability under regulatory auditability standards, we examine an active 15-node transaction cluster extracted from confirmed illicit transactions on Bitcoin Elliptic (Fig. 4). Saliency weights from GNNExplainer [54] show: (1) *Dispersal* ($S_0 \rightarrow M_{1...5}$): S_0 distributes \$48,500 across five




Fig. 4. Qualitative Attribution: 15-node Bitcoin Elliptic subgraph via GNNExplainer [54]. Dark red: illicit source (S_0); red: dispersal (M_1 – M_5); orange: layering mules (L_1 – L_5) with wash loops; green: exit accounts (C_1 – C_4); dashed gray: pruned chaff ($g_{ij} < 0.10$). Motifs, not legal adjudications.

accounts in sub-\$10,000 bursts over 22 minutes (structuring under 31 U.S.C. 5324); (2) *Layering* ($M \rightarrow L_{1...5}$): Intermediaries exhibit near-unity flow conservation ($\Phi_{\text{flow}} \approx 0.994$) and wash loops ($L_1 \rightarrow L_2 \rightarrow L_3 \rightarrow L_1$); (3) *Consolidation* ($L \rightarrow C_{1...4}$): Funds aggregate into four high-velocity exchange exit accounts; and (4) *Camouflage Filtration*: Edges to commercial nodes receive low gating ($g_{ij} \leq 0.042$), isolating the illicit core. Formally, *Factual Explanation Fidelity* is $\text{Fidelity}(G, G_s) = \hat{p}(y = 1 \mid G) - \hat{p}(y = 1 \mid G \setminus G_s)$. Pruning G_s drops illicit confidence from $\hat{p} = 0.982$ to $\hat{p} = 0.040$ (Fidelity = 0.942), with 96.40% motif precision across $N = 500$ clusters (Supplementary Table S6).

B. Cold-Start Boundary Case & Counterfactual Explanations: We analyze a cold-start mixer stub in `ibm_amlsim_li`: an account with no history ($\deg(u) = 0$) receives \$9,400 at $t_{\text{pred}} = t_{\text{deposit}}$. Standalone GNNs collapse ($\hat{p} = 0.28 < \tau = 0.50$), but C-STGB generates ambiguous set $\Gamma(X) = \{0, 1\}$ and routes to Tier 2 (Review Queue), averting an undetected false negative. Counterfactual analysis confirms normalizing amounts reduces attribution by 48.5% (FIN-2019-A003), while temporal dilation lowers Hawkes intensity below μ_0 (−38.2% risk). *Operational Recourse*: False-positive merchant holds link to an auditable checklist, allowing examiners to verify submissions and release holds without formal SAR escalation.

C. Auxiliary Decision-Support Prototype (Multi-Agent SAR Drafting): To bridge detections into compliance workflows under SR 26-2 [53], C-STGB incorporates an **Auxiliary Decision-Support Prototype** (Supplementary Fig. S4) decoupled from the core GNN. The pipeline coordinates: (1) an *Auditor Agent* checking OFAC sanctions; (2) an *Investigator Agent* extracting 2-hop subgraphs; and (3) a *Drafter Agent* compiling XML dossiers under FinCEN Form 111 with SHA-256 Merkle audit receipts. Autonomous filing is prohibited; filing authority resides exclusively with a certified BSA Officer. In a pilot evaluation across $N = 200$ audit dossiers (100 manual, 100 prototype-assisted), drafts achieved 99.2% statutory concordance ($\kappa = 1.0$ typology, $\kappa = 0.98$ factual grounding) in 28.4 s per case (vs. 45.0 min manual baseline).

VII. LIMITATIONS, THREATS TO VALIDITY & ETHICAL CONSIDERATIONS

1) Observability & Banking Silos: Fiat banking operates under statutory secrecy (RFP, GDPR), restricting institutions to local 1-hop visibility and motivating federated graph architectures (15.8%–44.5% F1 on IBM-AMLSim).

2) Ground-Truth Boundaries: High scores on crypto benchmarks reflect law enforcement seizures forming dense homophilic clusters ($h > 0.92$), validating algorithmic robustness

rather than unconstrained laundering evasion (e.g., Hawala or TBML).

3) Label Latency & Calibration: Illicit labels incur confirmation delays ($\Delta T_{\text{delay}} \in [30, 90]$ days). Assembling calibration buffers ($n_1 \geq 100$) requires pooling historical alerts; conformal bounds do not confer statutory clearance.

4) Cold-Start Invariants: On cold-start entities ($\deg(u) \leq 2$), C-STGB attenuates graph weights ($\alpha_u \rightarrow 0.08$) and transfers mass to flow invariants ($\mathbf{z}_{\text{inv}}$), securing 89.60% F1 via tabular velocity rather than graph topology.

5) Construct Overlap & Fairness: Edge-trust supervision uses hub thresholds ($\deg_{\text{out}} \geq 50$) on $\mathcal{E}_{\text{train}}$, sharing inductive bias with degree-matched camouflage; generalization is verified under gradient attacks. Demographic fairness warrants audits under ECOA.

6) Model Governance: LLM-assisted SAR drafting is strictly an auxiliary prototype; autonomous filings are prohibited under BSA regulations and Federal Reserve SR 26-2 [53]. Evaluated data comprise de-identified public ledgers or synthetic simulations with zero PII.

VIII. CONCLUSION & FUTURE DIRECTIONS

We presented **C-STGB**, an integrated geometric deep learning framework for anti-money laundering under extreme class skew, velocity evasion, and topological camouflage. C-STGB unifies continuous harmonic temporal attention, learnable edge-trust gating, typology-clustered latent GraphSMOTE, evidence-adaptive graph-tabular fusion, and Class-Conditional Conformal Risk Control. Across 14 benchmark networks (9.53M entities, 9.53M transactions across 182 evaluation pairs), C-STGB achieves 67.26% macro F1 (0.6848 PR-AUC), demonstrating statistically significant superiority over deep GNNs ($W = 105.0$, $p_{\text{adj}} < 0.001$) and competitive parity with tuned tree ensembles, alongside 0.45 ms Fast-Path streaming latency. An auxiliary prototype assists examiners with verifiable FinCEN Form 111 XML dossiers under human oversight (SR 26-2 [53]). Future work includes federated learning across banking silos, continual learning under drift, and hyperbolic representations.

ACKNOWLEDGMENT & AI DISCLOSURE

We thank Dept. of CSE, College of Technology, National University, Bangladesh, for computing facilities. Under IEEE GenAI Policy, AI tools were used solely for grammatical editing; an LLM was investigated strictly as an experimental object in Section VI-C. All formulations, proofs, and analyses were independently developed by the authors. Source code: <https://github.com/NazmulHasanNihal/Intelligent-AML>.

REFERENCES

- [1] United Nations Office on Drugs and Crime, “Estimating illicit financial flows resulting from drug trafficking and other transnational organized crimes,” UNODC, Vienna, Austria, Tech. Rep., 2011.
- [2] Financial Action Task Force, “Risk-based approach guidance for the banking sector,” FATF/OECD, Paris, France, Tech. Rep., 2021.
- [3] Z. Hu, Y. Dong, K. Wang, and Y. Sun, “Heterogeneous graph transformer,” in *Proceedings of The Web Conference (WWW)*, 2020, pp. 2704–2710.
- [4] S. Brody, U. Alon, and E. Yahav, “How attentive are graph attention networks?” in *International Conference on Learning Representations (ICLR)*, 2022.
- [5] T. N. Kipf and M. Welling, “Semi-supervised classification with graph convolutional networks,” in *International Conference on Learning Representations (ICLR)*, 2017.
- [6] W. L. Hamilton, R. Ying, and J. Leskovec, “Inductive representation learning on large graphs,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017, pp. 1024–1034.
- [7] M. Weber, G. Domeniconi, J. Chen, K. D. Weidele, C. Bellei, T. Robinson, and C. Shen, “Anti-Money Laundering in Bitcoin: Experimenting with graph convolutional networks for financial forensics,” in *Proceed-*

- [8] A. Pareja, G. Domeniconi, J. Chen, T. Teng, T. Wang, K. Chai, M. Magdon-Ismael, and C. Shen, "EvolveGCN: Evolving graph convolutional networks for dynamic graphs," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 34, no. 04, 2020, pp. 5363–5370.
- [9] Y. Dou, Z. Liu, L. Sun, Y. Deng, H. Peng, and P. S. Yu, "Enhancing graph neural network-based fraud detectors against camouflaged fraudsters," in *Proceedings of the 29th ACM International Conference on Information & Knowledge Management (CIKM)*, 2020, pp. 315–324.
- [10] C. Bellei, M. Xu, R. Phillips, T. Robinson, M. Weber, T. Kaler, C. E. Leiserson, Arvind, and J. Chen, "The shape of money laundering: Subgraph representation learning on the blockchain with the Elliptic2 dataset," *arXiv preprint arXiv:2404.19109*, 2024.
- [11] F. Johannessen and M. Jullum, "Finding money launderers using heterogeneous graph neural networks," *The Journal of Finance and Data Science*, vol. 11, p. 100175, 2025.
- [12] D. Cheng, Y. Ye, S. Xiang, Z. Ma, Y. Zhang, and C. Jiang, "Anti-money laundering by group-aware deep graph learning," *IEEE Transactions on Knowledge and Data Engineering*, vol. 35, no. 12, pp. 12444–12457, 2023.
- [13] Y. Fu, F. Liu, J. Zhang, P. S. Yu, and C. Shi, "Nowhere to H²IDE: Fraud detection from multi-relation graphs via disentangled homophily and heterophily identification," *IEEE Transactions on Knowledge and Data Engineering*, vol. 37, no. 3, pp. 1421–1435, 2025.
- [14] X. Qiao, Y. Zhao, X. Ma, Y. Zhu, C. Aggarwal, and J. Pei, "Deep graph anomaly detection: A survey and new perspectives," *IEEE Transactions on Knowledge and Data Engineering*, vol. 37, no. 9, pp. 5106–5126, 2025.
- [15] J. Wu, Q. Yuan, D. Lin, W. You, W. Chen, C. Chen, and Z. Zheng, "Who are the phishers? phishing scam detection on Ethereum via network embedding," *IEEE Transactions on Systems, Man, and Cybernetics: Systems*, vol. 52, no. 2, pp. 1156–1166, 2022.
- [16] P. Zheng, H.-N. Dai, J. Wu, and Z. Zheng, "XBlock-ETH: Extracting and analyzing blockchain data from Ethereum," *IEEE Open Journal of the Computer Society*, vol. 1, pp. 95–106, 2020.
- [17] N. Gandai, J. T. Hamrick, T. Moore, and T. Oberman, "Price manipulation in the Bitcoin ecosystem," *Journal of Monetary Economics*, vol. 95, pp. 86–96, 2018.
- [18] J. Lorenz, M. I. Silva, D. Aparício, J. T. Ascensão, and P. Bizarro, "Machine learning applied to anomaly detection in financial transactions: A review and an application to a multi-bank dataset," *arXiv preprint arXiv:2010.05779*, 2020.
- [19] T. Suzumura, Y. Zhou, N. Baracaldo, S. Guan, K. Hou, R. Houlihan, M. Stamatatos, P. Sapkota, and C. Zhang, "Towards online Anti-Money Laundering detection on distributed dynamic graphs," in *Proceedings of the IEEE International Conference on Big Data (Big Data)*, 2019, pp. 2350–2359.
- [20] A. Dal Pozzolo, G. Boracchi, O. Caelen, C. Alippi, and G. Bontempi, "Credit card fraud detection: A realistic modeling and a novel learning strategy," *IEEE Transactions on Neural Networks and Learning Systems*, vol. 29, no. 8, pp. 3784–3797, 2018.
- [21] C. Wang, B. Chen, G. Li, and H. Wang, "Automated graph neural network search under federated learning framework," *IEEE Transactions on Knowledge and Data Engineering*, vol. 35, no. 10, pp. 9959–9972, 2023.
- [22] E. Rossi, B. Zhou, F. Monti, F. Frasca, U. Alon, and M. M. Bronstein, "Temporal graph networks for deep learning on dynamic graphs," in *NeurIPS Workshop on Graph Representation Learning and Beyond*, 2020.
- [23] D. Xu, C. Ruan, E. Korpeoglu, S. Kumar, and K. Achan, "Inductive representation learning on temporal graphs," in *International Conference on Learning Representations (ICLR)*, 2020.
- [24] Y. Wang, K. Ding, C. Zhou, J. Liu, and S. Parthasarathy, "Dynamic graph anomaly detection with multi-scale temporal graph neural networks," in *Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*, 2024, pp. 3210–3221.
- [25] A. G. Hawkes, "Spectra of some self-exciting and mutually exciting point processes," *Biometrika*, vol. 58, no. 1, pp. 83–90, 1971.
- [26] H. Liu, Y. Yang, and X. Wang, "Overcoming catastrophic forgetting in graph neural networks," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 35, no. 10, 2021, pp. 8653–8661.
- [27] D. Zügner, A. Akbar Kazemi, and S. Günnemann, "Adversarial attacks on neural networks for graph data," in *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2018, pp. 2847–2856.
- [28] A. Bojchevski and S. Günnemann, "Adversarial attacks on node embeddings via graph poisoning," in *Proceedings of the 36th International Conference on Machine Learning (ICML)*, 2019, pp. 695–704.
- [29] R. Ou, K. Zhu, N. Zhang, M. Zhou, and C. Jiang, "CamFD: Semi-supervised camouflage-aware fraud detection based on dynamic graphs," *IEEE Transactions on Systems, Man, and Cybernetics: Systems*, vol. 56, no. 1, pp. 112–125, 2026.
- [30] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, "SMOTE: Synthetic minority over-sampling technique," *Journal of Artificial Intelligence Research*, vol. 16, pp. 321–357, 2002.
- [31] J. Davis and M. Goadrich, "The relationship between Precision-Recall and ROC curves," in *Proceedings of the 23rd International Conference on Machine Learning (ICML)*, 2006, pp. 233–240.
- [32] T. Zhao, X. Zhang, and S. Wang, "GraphSMOTE: Imbalanced node classification on graphs with graph neural networks," in *Proceedings of the 14th ACM International Conference on Web Search and Data Mining (WSDM)*, 2021, pp. 833–841.
- [33] V. Vovk, A. Gahmmerman, and G. Shafer, *Algorithmic Learning in a Random World*. Springer Science & Business Media, 2005.
- [34] A. N. Angelopoulos and S. Bates, "A gentle introduction to conformal prediction and distribution-free uncertainty quantification," *Foundations and Trends® in Machine Learning*, vol. 16, no. 4, pp. 494–591, 2023.
- [35] S. Bates, A. N. Angelopoulos, E. J. Candès, and M. I. Jordan, "Distribution-free risk-controlling prediction sets," *Journal of the ACM*, vol. 68, no. 6, pp. 1–34, 2021.
- [36] K. Huang, Y. Jin, E. J. Candès, and J. Leskovec, "Uncertainty quantification over graph neural networks via conformal prediction," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 46, no. 11, pp. 8120–8135, 2024.
- [37] T. Ding, A. N. Angelopoulos, S. Bates, M. I. Jordan, and R. J. Tibshirani, "Class-conditional conformal prediction with many classes," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 36, 2023, pp. 4461–4475.
- [38] Y. Romano, M. Sesia, and E. J. Candès, "Classification with valid and adaptive coverage," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, 2020, pp. 3581–3591.
- [39] R. J. Tibshirani, E. J. Candès, A. Ramdas, and R. Foygel Barber, "Conformal prediction under covariate shift," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.
- [40] R. F. Barber, E. J. Candès, A. Ramdas, and R. J. Tibshirani, "Conformal prediction beyond exchangeability," *The Annals of Statistics*, vol. 51, no. 2, pp. 816–845, 2023.
- [41] I. Gibbs and E. J. Candès, "Adaptive conformal inference under distribution shift," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, 2021, pp. 1660–1672.
- [42] C. Zhang, Y. Wang, Z. Liu, C. Wang, and J. Li, "Graph neural networks for financial fraud detection: A comprehensive review," *Frontiers of Computer Science*, vol. 19, no. 3, p. 193801, 2025.
- [43] T. Pourhabibi, K.-L. Ong, B. H. Kam, and Y. L. Boo, "Graph learning in financial fraud detection: A systematic survey on dynamic graphs, robustness, and future directions," *Journal of Financial Crime and Data Science*, vol. 2, no. 1, p. 100033, 2025.
- [44] W. Chen, H. Liu, Z. Zhao, M. Zhang, and L. Xu, "Uncertainty-aware spatio-temporal graph neural architecture for financial fraud detection under relational volatility," *Scientific Reports*, vol. 16, no. 1, p. 55651, 2026.
- [45] E. A. Lopez-Rojas and S. Axelsson, "PaySim: A financial mobile money simulator for fraud detection," in *Proceedings of the 28th European Modeling and Simulation Symposium (EMSS)*, 2016, pp. 249–255.
- [46] X. Huang, Y. Shen, Y. Huang, Z. Zhang, X. Zhou, X. Zhang, Y. Du, and B. Hu, "DGraph: A large-scale financial dataset for graph anomaly detection," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 35, 2022, pp. 22765–22777.
- [47] A. Dal Pozzolo, O. Caelen, R. A. Johnson, and G. Bontempi, "Calibrating probability with undersampling for unbalanced classification," in *IEEE Symposium Series on Computational Intelligence (SSCI)*, 2015, pp. 159–166.
- [48] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 785–794.
- [49] L. Prokhorenkova, G. Gusev, A. Vorobev, A. V. Dorogush, and A. Gulin, "CatBoost: Unbiased boosting with categorical features," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 31, 2018.
- [50] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, "LightGBM: A highly efficient gradient boosting decision tree," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017, pp. 3146–3154.
- [51] P. Velićković, G. Cucurull, A. Casanova, A. Romero, P. Liò, and Y. Bengio, "Graph attention networks," in *International Conference on Learning Representations (ICLR)*, 2018.
- [52] K. Xu, W. Hu, J. Leskovec, and S. Jegelka, "How powerful are graph neural networks?" in *International Conference on Learning Representations (ICLR)*, 2019.
- [53] Federal Reserve Board, "Supervisory letter sr 26-2: Guidance on model risk management for advanced machine learning," FRB, OCC, FDIC, Washington, DC, Tech. Rep., 2026.
- [54] R. Ying, D. Bourgeois, J. You, M. Zitnik, and J. Leskovec, "GNNExplainer: Generating explanations for graph neural networks," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.

Md. Nazmul is pursuing the B.Sc. (Hons.) degree in CSE at College of Technology, National University, Bangladesh. Research focuses on graph representation learning, spatio-temporal neural networks, and conformal risk control.

Musrat Jahan Gungun is pursuing the B.Sc. (Hons.) degree in CSE at College of Technology, National University, Bangladesh. Research focuses on machine learning, algorithmic fairness, and topological anomaly detection.

Maheli Ahmed is a Faculty Member in CSE at College of Technology, National University, Bangladesh. Research interests include machine learning, deep learning, model governance, and fraud surveillance.