

Supplementary Material: Risk-Controlled Spatio-Temporal Graph Learning for Anti-Money Laundering Under Extreme Imbalance and Topological Camouflage

Md. Nazmul, Musrat Jahan Gungun, and Maheli Ahmed

TABLE S1
SUMMARY OF MATHEMATICAL NOTATIONS, TENSOR DIMENSIONS, AND OPERATORS.

Symbol	Mathematical Definition	Dimension / Type
$\mathcal{G} = (\mathcal{V}, \mathcal{E})$	Continuous-time dynamic transaction graph	Graph object
$\mathbf{X} \in \mathbb{R}^{ \mathcal{V} \times d_v}$	Node input attribute matrix	Real matrix
$\mathbf{E} \in \mathbb{R}^{ \mathcal{E} \times d_e}$	Transaction edge feature matrix	Real matrix
$\Delta_{t_{uv}}$	Continuous time delta ($t_v - t_u \geq 0$)	Non-negative scalar
$\Phi_{\text{time}}(\Delta t)$	Harmonic continuous temporal positional embedding	$\mathbb{R}^{d_t}$ vector
$g_{ij} \in (0, 1)$	Context-aware adaptive edge trust gate	Scalar gate
$\Phi_{\text{flow}}(u, t)$	Mass flow conservation ratio	Scalar $\in [0, \infty)$
$\mathcal{A}_u(t)$	Hawkes process self-exciting transaction intensity	Non-negative scalar
$\mathbf{s}_{\text{fwd}}, \mathbf{s}_{\text{bwd}}$	Temporally restricted forward and reverse PageRank taint vectors	$\mathbb{R}^{ \mathcal{V} }$ vectors
$\mathbf{z}_{\text{inv}}(u, t)$	12-dimensional canonical representation	$\mathbb{R}^{12}$ vector
$\mathbf{h}_u^{(l)}$	Hidden state representation of node u at layer l	$\mathbb{R}^d$ vector
C_k	Latent typology cluster of illicit transactions	Subset $C_k \subset \mathcal{V}_{\text{illicit}}$
$\alpha_u \in [0, 1]$	Evidence-adaptive graph-tabular fusion weight	Scalar weight
$\mathbf{h}_u^a$	Unified anomaly representation vector	$\mathbb{R}^{d'}$ vector
τ^*	Cost-sensitive Bayes risk decision threshold	Optimal scalar $\in (0, 1)$
$\hat{q}^{(y)}$	Class-conditional conformal quantile ($y \in \{0, 1\}$)	Calibration scalar
$\Gamma(X)$	Finite-sample conformal prediction set	Discrete subset $\subseteq \{0, 1\}$
α	Target significance error level ($1 - \alpha$ coverage)	Scalar $\in (0, 1)$

APPENDIX A

SUMMARY OF MATHEMATICAL NOTATIONS AND OPERATORS

Table S1 provides the comprehensive reference glossary of mathematical symbols, tensor dimensions, and operator definitions employed throughout C-STGB.

APPENDIX B

COMPLIANCE RULE VERIFIER AST GRAMMAR & SCHEMA SPECIFICATION

To guarantee regulatory determinism and enforce zero entity hallucinations during assisted Suspicious Activity Report (SAR) narrative draft generation (Section VI of the main paper), the multi-agent forensic swarm employs a formal Abstract Syntax Tree (AST) grammar and runtime Pydantic schema validator.

A. Extended Backus-Naur Form (EBNF) Grammar

All natural-language and structured FinCEN Form 111 XML narratives generated by the SAR Drafter Agent must strictly conform to the following formal grammar $\mathcal{G}_{\text{SAR}}$:

Md. Nazmul and Musrat Jahan Gungun are with the Department of Computer Science and Engineering, College of Technology, National University, Gazipur 1704, Bangladesh (e-mail: nazmulhas36@gmail.com, gungunjanha84@gmail.com).

Maheli Ahmed (Faculty Supervisor) is with the Department of Computer Science and Engineering, College of Technology, National University, Gazipur 1704, Bangladesh (e-mail: maheli.ahmed.cse.cot@gmail.com).

Corresponding author: Md. Nazmul.

```

SAR_Document ::= Header
                SubjectSection
                ActivitySection
                EvidenceSection
                AuditTrailer

Header ::= "<SAR_Form111"
        " Version='2.0'>"
        FilingInst
        FilingTimestamp
        CaseID
        "</SAR_Form111>"

SubjectSection ::= "<Subjects>"
                Subject+
                "</Subjects>"

Subject ::= "<Subject ID="
        EntityID ">"
        IdentityType
        RiskScore
        ConformalTier
        Jurisdiction
        "</Subject>"

ActivitySec ::= "<Activity>"
                Typology
                Amount TemporalSpan
                Narrative
                "</Activity>"

Typology ::= "Smurfing"
        | "CircularWash"
        | "RapidLayering"
        | "DormantPass"
        | "HubCamouflage"

EvidenceSec ::= "<Evidence>"
                Root CausalHop+
                FlowRatio Hawkes
                "</Evidence>"

CausalHop ::= "<Hop Seq=" Int ">"
                Src "-->" Dst
                Amt Time Gate
                "</Hop>"

FlowRatio ::= "<Flow In=" Real
        " Out=" Real
        " Ratio=" Real "/>"

AuditTrailer ::= "<Audit_Governance_SR26_2>"
                MerkleRoot
                SHA256Proof Time
                "</Audit_Governance_SR26_2>"

EntityID ::= [a-zA-Z0-9_]{1,64}
Real ::= [0-9]+ "." [0-9]+
Int ::= [0-9]+

```

B. Two-Layer Runtime Schema & Factual Graph Validation

Before any research SAR payload is accepted into the audit repository or dispatched to compliance officers (Tier 2), it is compiled into an AST and subjected to two deterministic verification layers: (1) syntactic Pydantic schema validation $\mathcal{S}_{\text{AML}}$, and (2) factual graph membership assertion against the causal subgraph $\mathcal{G}_{\text{sub}} = (\mathcal{V}_{\text{sub}}, \mathcal{E}_{\text{sub}})$.

```

class CausalHopSchema(BaseModel):
    hop_seq: int = Field(..., ge=1, le=10)
    src_id: str = Field(..., min_length=1)

```

TABLE S2

SYSTEMATIC METHODOLOGICAL COMPARISON: STRUCTURAL AND OPERATIONAL CAPABILITIES ACROSS MODEL PARADIGMS IN FRAUD AND AML SURVEILLANCE.

Model / Paradigm Architecture	Continuous Temporal Dynamics Representation	Topological Camouflage Edge-Trust Gating	Latent Typology Graph Oversampling	Cold-Start Graph-Tabular Fusion	Finite-Sample Conformal Coverage	Streaming Single Inference Latency
XGBoost / CatBoost / LightGBM [1]–[3]	Feature-level Δt	N/A (Tabular only)	Euclidean SMOTE	N/A (Tabular only)	Heuristic Score τ	< 0.20 ms
GCN / GAT / GIN [4]–[6]	Static Snapshot	Uniform / Homophilic	None	Zero-padded features	Softmax probability	> 12.5 ms
EvolveGCN [7]	Discrete time slices	Unweighted edge steps	None	Isolated nodes	Uncalibrated	> 45.0 ms
TGN / TGAT [8], [9]	Single decay $\exp(-\lambda \Delta t)$	Unfiltered edges	None	Node memory state	Uncalibrated	> 8.5 ms
CARE-GNN (Adversarial) [10]	Static graph	RL edge selection	None	Concatenation	Uncalibrated	> 18.0 ms
GraphSMOTE [11]	Static graph	Unfiltered edges	Global unconstrained	No fusion fallback	Uncalibrated	> 14.2 ms
Johannessen & Jullum [12]	Static heterogeneous	Static aggregation	None	Feature concatenation	Uncalibrated	> 16.5 ms
Wang <i>et al.</i> (ACM KDD, 2024) [13]	Continuous multi-scale	Unfiltered edges	None	No invariant fallback	Uncalibrated	> 5.8 ms
C-STGB (Proposed Framework)	Tri-band harmonic Φ_{time}	Learnable gate $\hat{g}_{ij}$	Typology manifold C_k	Evidence-adaptive α_u	Class-conditional $\geq 99\%$	0.45 ms Fast-Path

TABLE S3

COMPONENT-LEVEL NOVELTY AND ARCHITECTURAL DIFFERENTIATION OF C-STGB RELATIVE TO CLOSEST PRIOR ART.

Framework Component	Closest Prior Work	What Prior Method Does	What C-STGB Modifies	New Theoretical / Empirical Property
Temporal Attention	TGN [8], TGAT [9]	Single exponential decay $\exp(-\lambda \Delta t)$	Continuous tri-band harmonic attention Φ_{time}	Unifies sub-second velocity bursts with multi-week dormancy
Edge-Trust Gating	CARE-GNN [10]	RL-based relation selection on static graphs	Multi-modal MLP gate $\hat{g}_{ij}$ with temporal features	Filters camouflage chaff while mitigating commercial hub over-smoothing
Imbalance Oversampling	GraphSMOTE [11]	Unconstrained global latent interpolation	Typology-clustered interpolation along manifolds C_k	Preserves point-level flow regularity [min, max] without distortion
Graph-Tabular Fusion	Dual-Stream Stacking	Static concatenation or linear stacking	Evidence-adaptive gating parameter α_u	Dynamic fallback to flow descriptors on cold-start stubs ($d \leq 2$)
Uncertainty Triage	Split Conformal [14], [15]	Marginal transductive coverage on static data	Class-conditional CRC with streaming ACI updates	Guaranteed $1 - \alpha$ class-conditional coverage under temporal shift

```

dst_id: str = Field(..., min_length=1)
amt_usd: float = Field(..., gt=0.0)
timestamp_iso: str
edge_gate: float = Field(..., ge=0, le=1)

class SARDocumentSchema(BaseModel):
    case_id: str = Field(...,
        regex=r"SAR-[0-9]{4}-[A-Z0-9]{8}$")
    filing_ts: datetime
    tier: Literal["Tier_1", "Tier_2"]
    risk_score: float = Field(...,
        ge=0.0, le=1.0)
    typology: Literal[
        "Smurfing", "CircularWash",
        "RapidLayering", "DormantPass",
        "HubCamouflage"
    ]
    flow_ratio: float
    trajectory: List[CausalHopSchema] = \
        Field(..., min_items=1)
    merkle_hash: str = Field(...,
        min_length=64, max_length=64)

    @validator("trajectory")
    def verify_topological_chain(cls, v):
        for i in range(len(v) - 1):
            assert v[i].dst_id == \
                v[i+1].src_id, \
                "Disconnected chain."
        return v

def verify_factual_graph_membership(
    doc: SARDocumentSchema,
    V_sub: set, E_sub: set
) -> bool:
    """Layer 2: Factual graph assertion."""
    for hop in doc.trajectory:
        assert hop.src_id in V_sub, \
            f"Node hallucination: {hop.src_id}"
        assert hop.dst_id in V_sub, \
            f"Node hallucination: {hop.dst_id}"
        assert (hop.src_id,
            hop.dst_id) in E_sub, \
            f"Edge hallucination: {hop.src_id}"
    return True

```

C. Unsupported-Entity Prevention Invariant & Runtime Compiler Specification

Proposition S1 (Deterministic Unsupported-Entity Prevention Invariant (Proposition S1)). *Let $\mathcal{G}_{\text{sub}} = (\mathcal{V}_{\text{sub}}, \mathcal{E}_{\text{sub}})$ denote the factual transaction subgraph extracted from the immutable ledger, and let $\mathcal{T}(\text{SAR})$ denote the parsed AST generated by the LLM-based SAR Drafter Agent. If $\mathcal{T}(\text{SAR})$ passes schema validation $\mathcal{S}_{\text{AML}}$ and deterministic factual graph validation*

verify_factual_graph_membership, then:

$$\forall u \in \mathcal{V}(\mathcal{T}), u \in \mathcal{V}_{\text{sub}} \quad \text{and} \quad \forall (u, v) \in \mathcal{E}(\mathcal{T}), (u, v) \in \mathcal{E}_{\text{sub}} \quad (1)$$

Proof. Layer 1 enforces syntactical AST typing and sequence continuity. Layer 2 executes deterministic set-membership assertions over the factual sets $\mathcal{V}_{\text{sub}}$ and $\mathcal{E}_{\text{sub}}$. Any entity identifier or edge generated in the LLM narrative that does not strictly map to an entry in $\mathcal{V}_{\text{sub}}$ or $\mathcal{E}_{\text{sub}}$ raises an immediate runtime `AssertionError`, triggering a constrained retry under greedy decoding ($T = 0.0$). Thus, unsupported entity references and hallucinated relations cannot enter the finalized compliance record. $\square$

Remark (Compiler-Enforced Invariants & Human-in-the-Loop Sign-Off): We emphasize that Proposition S1 represents a deterministic software compiler and AST schema-enforced verification invariant rather than an analytical mathematical guarantee inherent to unconstrained autoregressive neural text generators. In production compliance operations, machine verification serves strictly as an operational pre-filter; final filing approval requires manual inspection and digital cryptographic countersignature by a certified Bank Secrecy Act (BSA) Officer under Federal Reserve SR 26-2 model governance guidelines [16].

APPENDIX C

ALGORITHMIC PROCEDURES & PIPELINE ARCHITECTURE

This section formalizes the multi-stage training routine (Algorithm S1) and streaming online inference triage engine (Algorithm S2) corresponding to the end-to-end execution flowchart detailed in Fig. S1.

APPENDIX D

EXTENDED MATHEMATICAL PROOFS & CONFORMAL DERIVATIONS

A. Dvoretzky-Kiefer-Wolfowitz (DKW) Finite-Sample Bound under TV Drift

In continuous-time streaming financial networks, classical exchangeability is challenged by continuous-time temporal

Unified Multi-Layer C-STGB End-to-End Pipeline Execution Flow (Algorithms S1 & S2)




Fig. S1. Unified Multi-Layer C-STGB End-to-End Pipeline Execution Flow (Algorithms S1 and S2), detailing data flow through temporal encoding, edge trust filtering, adaptive fusion, risk optimization, and 3-tier selective abstention triage.

Algorithm S1 C-STGB Unified Multi-Stage Training Pipeline

Require: Graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$, Features $\mathbf{X}$, Timestamps t , Target α , Weights $\gamma_1 = 0.50, \gamma_2 = 0.10$.

Ensure: Trained model parameters Θ , Calibrated threshold τ^* , Conformal quantiles $\hat{q}^{(0)}, \hat{q}^{(1)}$.

- 1: Compute 12-D canonical invariants $\mathbf{z}_{\text{inv}}(u)$ for all $u \in \mathcal{V}$.
- 2: Cluster minority nodes into latent typologies $C_1, \dots, C_K$.
- 3: Synthesize minority representations $\mathbf{h}_{\text{syn}}$ and bilinear edges $\hat{\mathbf{A}}_{\text{syn}}$.
- 4: Mine hard-negative commercial hubs into batch bank $\mathcal{B}_{\text{hard}}$.
- 5: **for** epoch = 1 to N_{epochs} **do**
- 6: Compute multi-scale temporal embeddings $\Phi_{\text{time}}(\Delta t_{ij})$ and weights $w_{ij}(t)$.
- 7: Evaluate context-aware edge trust gates $\mathbf{g}_{ij}$ and filter camouflage links.
- 8: Aggregate degree-capped temporal messages to obtain $\mathbf{h}_u^{(L)}$.
- 9: Compute evidence-adaptive fusion weights α_u and representations $\mathbf{h}_u^*$.
- 10: Evaluate loss $\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{task}} + \gamma_1 \mathcal{L}_{\text{recon}} + \gamma_2 \mathcal{L}_{\text{aux}}$.
- 11: Update Θ via AdamW with cosine learning rate annealing.
- 12: **end for**
- 13: Optimize Bayes decision threshold τ^* on validation split $\mathcal{D}_{\text{val}}$.
- 14: Calibrate class-conditional non-conformity quantiles $\hat{q}^{(0)}, \hat{q}^{(1)}$ on $\mathcal{D}_{\text{cal}}$.
- 15: **Return** $\Theta, \tau^*, \hat{q}^{(0)}, \hat{q}^{(1)}$.

autocorrelation and non-stationary behavioral drift. As formulated in Section III-I of the main manuscript [17]–[19], C-STGB employs online Adaptive Conformal Inference (ACI) tracking and a rolling calibration split $\mathcal{D}_{\text{cal}}(t)$ over window W_{cal} as the explicit mechanisms bounding total variation (d_{TV}) error.

Definition S1 (Calibration Holdout Abstraction and Mixing Assumption). To apply empirical process concentration to the

Algorithm S2 Streaming Online Inference & Risk-Controlled Triage

Require: Streaming transaction $(u, v, t, \mathbf{e}_{uv})$, Parameters $\Theta, \tau^*, \hat{q}^{(0)}, \hat{q}^{(1)}$.

Ensure: Operational Triage Action $\in \{\text{Auto-Block, Review Queue, Clear}\}$.

- 1: Ingest transaction event and update temporally restricted ego-graph ($t_e \leq t$).
- 2: Compute canonical invariants $\mathbf{z}_{\text{inv}}(u)$ and edge trust gate $\mathbf{g}_{uv}$.
- 3: Extract temporal message embeddings $\mathbf{h}_u^{(L)}$.
- 4: Compute adaptive fusion weight α_u and fused representation $\mathbf{h}_u^*$.
- 5: Predict anomaly probability $\hat{p}_u = \sigma(\text{MLP}(\mathbf{h}_u^*))$.
- 6: Construct class-conditional conformal prediction set $\Gamma(u) \subseteq \{0, 1\}$.
- 7: **if** $\Gamma(u) == \{1\}$ **then**
- 8: **Route to Tier 1:** Quarantined Asset Hold & Compliance Decision Support Dossier.
- 9: **else if** $\Gamma(u) == \{0, 1\}$ **then**
- 10: **Route to Tier 2:** Selective Abstention $\perp \rightarrow$ Compliance Officer Queue.
- 11: **else**
- 12: **Route to Tier 3:** Straight-Through Clear.
- 13: **end if**

calibration split, we assume that calibration non-conformity scores $\mathcal{D}_{\text{cal}}^{(y)}$ are drawn from a stationary holdout window under bounded temporal α -mixing: $\alpha(k) \leq C \exp(-ck)$ for event separation k , such that inter-alert draws beyond temporal decorrelation span τ_{mix} satisfy the exchangeable calibration abstraction. When drawing independent alert cases from historical audit pools, the calibration split strictly satisfies within-class exchangeability.

Definition S2 (Total Variation Distance and Kolmogorov-Smirnov Metric). For any two Borel probability measures $\mathcal{P}, \mathcal{Q}$ on $(\mathbb{R}, \mathcal{B})$, the Total Variation distance is defined as $d_{\text{TV}}(\mathcal{P}, \mathcal{Q}) = \sup_{A \in \mathcal{B}} |\mathcal{P}(A) - \mathcal{Q}(A)|$. The Kolmogorov-

Smirnov distance is defined as $d_{KS}(\mathcal{P}, \mathcal{Q}) = \sup_{s \in \mathbb{R}} |F_{\mathcal{P}}(s) - F_{\mathcal{Q}}(s)|$, which satisfies $d_{KS}(\mathcal{P}, \mathcal{Q}) \leq d_{TV}(\mathcal{P}, \mathcal{Q})$. In C-STGB, empirical two-sample KS tests over non-conformity scores provide a computable hypothesis test for detecting operational drift, while d_{TV} bounds the theoretical worst-case coverage error across all measurable decision sets.

Proposition S2 (Non-Asymptotic Conformal TV-Drift Error Bound). *Under locally bounded total variation drift $d_{TV}(\mathcal{P}_t^{(y)}, \mathcal{P}_{cal}^{(y)}) \leq \epsilon_{drift}$ and calibration mixing (Definition S1), the non-asymptotic class-conditional coverage error bound decomposes cleanly via maximal coupling into an environmental distribution shift component and an empirical process concentration component:*

$$\begin{aligned} & |\mathbb{P}_t(Y \in \Gamma(X) \mid Y = y) - (1 - \alpha_y)| \\ & \leq d_{TV}(\mathcal{P}_t^{(y)}, \mathcal{P}_{cal}^{(y)}) + \sup_s \left| F_{cal}^{(y)}(s) - \hat{F}_{cal}^{(y)}(s) \right| + \frac{1}{n(y)} \end{aligned} \quad (2)$$

where $n(y) = |\mathcal{D}_{cal}^{(y)}(t)|$ denotes the class-conditional calibration sample size.

Proof. Let $F_t^{(y)}(s) = \mathbb{P}_t(S^{(y)} \leq s)$ denote the true cumulative distribution function (CDF) of class-conditional non-conformity scores under the streaming test distribution $\mathcal{P}_t$, and let $F_{cal}^{(y)}(s) = \mathbb{P}_{cal}(S^{(y)} \leq s)$ denote the true CDF under the rolling calibration distribution $\mathcal{P}_{cal}^{(y)}$. Let $\hat{F}_{cal}^{(y)}(s) = \frac{1}{n(y)} \sum_{i=1}^{n(y)} \mathbb{I}(s_i^{(y)} \leq s)$ denote the empirical CDF computed over the calibration split $\mathcal{D}_{cal}(t)$, where $n(y) = |\mathcal{D}_{cal}^{(y)}(t)|$.

By definition of the conformal quantile $\hat{q}^{(y)}$, the test coverage probability evaluates to $\mathbb{P}_t(Y \in \Gamma(X) \mid Y = y) = \mathbb{P}_t(S^{(y)} \leq \hat{q}^{(y)}) = F_t^{(y)}(\hat{q}^{(y)})$. Applying the triangle inequality:

$$\begin{aligned} |F_t^{(y)}(\hat{q}^{(y)}) - (1 - \alpha_y)| & \leq |F_t^{(y)}(\hat{q}^{(y)}) - F_{cal}^{(y)}(\hat{q}^{(y)})| \\ & \quad + |F_{cal}^{(y)}(\hat{q}^{(y)}) - \hat{F}_{cal}^{(y)}(\hat{q}^{(y)})| \\ & \quad + |\hat{F}_{cal}^{(y)}(\hat{q}^{(y)}) - (1 - \alpha_y)| \end{aligned} \quad (3)$$

By the maximal coupling characterization of total variation distance, for any Borel set $A \subseteq \mathbb{R}$:

$$|\mathbb{P}_t(S^{(y)} \in A) - \mathbb{P}_{cal}(S^{(y)} \in A)| \leq d_{TV}(\mathcal{P}_t^{(y)}, \mathcal{P}_{cal}^{(y)}) \quad (4)$$

Setting $A = (-\infty, \hat{q}^{(y)}]$ yields $|F_t^{(y)}(\hat{q}^{(y)}) - F_{cal}^{(y)}(\hat{q}^{(y)})| \leq d_{TV}(\mathcal{P}_t^{(y)}, \mathcal{P}_{cal}^{(y)})$.

Next, applying the Dvoretzky-Kiefer-Wolfowitz (DKW) inequality (Massart, 1990) to the exchangeable calibration sample $\mathcal{D}_{cal}^{(y)}$ (Definition S1):

$$\mathbb{P} \left(\sup_{s \in \mathbb{R}} |F_{cal}^{(y)}(s) - \hat{F}_{cal}^{(y)}(s)| \leq \sqrt{\frac{\ln(2/\delta)}{2n(y)}} \mid Y = y \right) \geq 1 - \delta \quad (5)$$

Finally, by the definition of the discrete empirical quantile with sample size $n(y)$, the discretization residual satisfies $|\hat{F}_{cal}^{(y)}(\hat{q}^{(y)}) - (1 - \alpha_y)| \leq \frac{1}{n(y)}$. Under non-atomic score continuity, combining these terms establishes the non-asymptotic finite-sample bound with high probability $1 - \delta$.

Adaptive Conformal Inference (ACI) Convergence under Delayed Feedback: To track non-stationary drift without parametric priors, ACI updates the effective target $\alpha_t \leftarrow \alpha_{t-1} + \gamma(\alpha - \text{err}_{t-\tau(t)})$ over the confirmed observation queue

with step size $\gamma > 0$ and confirmation delay $\tau(t) \leq \Delta T_{\text{delay}}$. By the Azuma-Hoeffding inequality for delayed Martingale difference sequences [19], long-run average coverage satisfies $\lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T \text{err}_t = \alpha$ almost surely, with finite-horizon tracking error bounded by $O(\gamma \tau_{\max} + \sqrt{\frac{\ln(1/\delta)}{T}})$. In streaming deployment, maintaining separate trackers $\alpha_t^{(0)}, \alpha_t^{(1)}$ provides class-adaptive empirical tracking; we explicitly distinguish this empirical tracking heuristic from finite-sample class-conditional theorems. $\square$

B. Chronological Temporal Purity and Zero-Leakage of Latent GraphSMOTE

Lemma S1 (Chronological Isolation of Latent Topological Oversampling). *Let $\mathcal{D}_{train} = \mathcal{G}_{[0, t_{train}]}$, $\mathcal{D}_{val} = \mathcal{G}_{(t_{train}, t_{val}]}$, $\mathcal{D}_{cal} = \mathcal{G}_{(t_{val}, t_{cal}]}$, and $\mathcal{D}_{test} = \mathcal{G}_{(t_{cal}, t_{test}]}$ denote a strictly chronological 4-way temporal partition. Let $\mathbf{h}_{syn}$ be any synthetic minority representation generated via convex combination of seed nodes $u, v \in C_k \subseteq \mathcal{V}_{train}$ with birth timestamp $t_{syn} = \max(t_u, t_v) \leq t_{train}$. Under the causal indicator constraint $\hat{\mathbf{A}}_{syn, w} = \sigma(\mathbf{h}_{syn}^T \mathbf{W}_{edge} \mathbf{h}_w) \cdot \mathbb{I}(t_w \leq t_{syn})$, the synthesized graph $\tilde{\mathcal{G}}_{train}$ satisfies:*

$$\begin{aligned} \mathcal{V}(\tilde{\mathcal{G}}_{train}) \cap (\mathcal{V}_{val} \cup \mathcal{V}_{cal} \cup \mathcal{V}_{test}) &= \emptyset, \quad \text{and} \\ \forall (i, j) \in \mathcal{E}(\tilde{\mathcal{G}}_{train}), \quad \max(t_i, t_j) &\leq t_{train} \end{aligned} \quad (6)$$

Consequently, synthetic edge formation induces zero forward lookahead bias and zero backpropagation gradient flow into prospective evaluation partitions.

Proof. By construction, parent seed nodes $u, v \in \mathcal{V}_{train}$, which implies $t_u, t_v \leq t_{train}$ and hence $t_{syn} \leq t_{train}$. Candidate target nodes w are strictly restricted to $C_{cand}(\mathbf{h}_{syn}) \cap \mathcal{V}_{train}^{\leq t_{syn}}$. The causal indicator $\mathbb{I}(t_w \leq t_{syn})$ evaluates to 0 for any node with $t_w > t_{syn}$, enforcing $\hat{\mathbf{A}}_{syn, w} = 0$. Because $t_{syn} \leq t_{train} < t_{val} < t_{cal} < t_{test}$, no candidate $w \in \mathcal{V}_{val} \cup \mathcal{V}_{cal} \cup \mathcal{V}_{test}$ can satisfy $t_w \leq t_{syn}$. Furthermore, synthetic node representations $\mathbf{h}_{syn}$ and virtual edges are discarded prior to evaluation on $\mathcal{D}_{val}, \mathcal{D}_{cal}$, and $\mathcal{D}_{test}$. Hence, mutual information satisfies $I(\tilde{\mathcal{G}}_{train}; \mathcal{D}_{test} \mid \mathcal{D}_{train}) = 0$. $\square$

C. Kirchhoff Mass-Balance Flow Conservation Invariant for Layering Cycles

Lemma S2 (Conservation of Flow in Pass-Through Conduit Chains). *Let $\mathcal{P}_{chain} = (v_0, v_1, \dots, v_K)$ denote an ordered multi-hop fund layering chain with transaction amounts $A(v_{k-1}, v_k)$ and retention loss rate $\epsilon_k \in [0, 1]$ at each hop (representing service fees or split commissions). If intermediate nodes v_k ($1 \leq k < K$) act strictly as pass-through conduit accounts without capital injection, then the aggregate flow conservation invariant satisfies:*

$$\Phi_{flow}(v_k, t) = \frac{\sum_{e \in \mathcal{E}_{out}(v_k)} A(e)}{\sum_{e \in \mathcal{E}_{in}(v_k)} A(e) + \epsilon} \in [1 - \bar{\epsilon}, 1.0] \quad (7)$$

where $\bar{\epsilon} = \max_k \epsilon_k < 0.15$. Consequently, the 12-D canonical feature $\tilde{\Phi}_{flow}(v_k, t) = \log(1 + \Phi_{flow}(v_k, t))$ is strictly bounded within $[\log(2 - \bar{\epsilon}), \log(2)] \approx [0.615, 0.693]$, forming an invariant topological signature that is invariant to arbitrary graph scale and entity degree.

Proof. By the definition of intermediate pass-through mules in financial layering (BSA 31 CFR §1010.100), the net retained capital $R(v_k) = \sum_{\text{in}} A - \sum_{\text{out}} A$ is bounded exclusively by transaction clearing fees $\epsilon_k \sum_{\text{in}} A$. Thus $\sum_{\text{out}} A = (1 - \epsilon_k) \sum_{\text{in}} A$. Dividing by $\sum_{\text{in}} A$ yields $\Phi_{\text{flow}}(v_k, t) = 1 - \epsilon_k$. Since $\epsilon_k \in [0, \bar{\epsilon}]$, the ratio $\Phi_{\text{flow}}(v_k, t)$ lies in $[1 - \bar{\epsilon}, 1.0]$. Applying the monotonic logarithmic transformation $\tilde{\Phi}_{\text{flow}} = \log(1 + \Phi_{\text{flow}})$ preserves the interval order, yielding the bounded invariant range. $\square$

Proposition S3 (Convex Combination Flow Invariant Preservation (Proposition S2)). *Let $u, v \in C_k$ be two parent pass-through accounts in the same latent typology cluster with total inflows $I_u, I_v > 0$ and outflows $O_u = \Phi_u I_u, O_v = \Phi_v I_v$, where $\Phi_u, \Phi_v \in [1 - \bar{\epsilon}, 1.0]$. Let the synthetic representation $\mathbf{h}_{\text{syn}} = \lambda \mathbf{h}_u + (1 - \lambda) \mathbf{h}_v$ ($\lambda \in [0, 1]$) linearly interpolate parent flow amounts, such that $I_{\text{syn}} = \lambda I_u + (1 - \lambda) I_v$ and $O_{\text{syn}} = \lambda O_u + (1 - \lambda) O_v$. Then the synthetic flow conservation ratio $\Phi_{\text{flow}}(\mathbf{h}_{\text{syn}}) = \frac{O_{\text{syn}}}{I_{\text{syn}}}$ satisfies:*

$$\min(\Phi_u, \Phi_v) \leq \Phi_{\text{flow}}(\mathbf{h}_{\text{syn}}) \leq \max(\Phi_u, \Phi_v) \quad (8)$$

In particular, $\Phi_{\text{flow}}(\mathbf{h}_{\text{syn}}) \in [1 - \bar{\epsilon}, 1.0]$.

Proof. Without loss of generality, assume $\Phi_u \leq \Phi_v$. Since $\lambda \in [0, 1]$ and $I_u, I_v > 0$, we have:

$$\begin{aligned} O_{\text{syn}} &= \lambda \Phi_u I_u + (1 - \lambda) \Phi_v I_v \\ &\geq \lambda \Phi_u I_u + (1 - \lambda) \Phi_u I_v = \Phi_u I_{\text{syn}}, \\ O_{\text{syn}} &= \lambda \Phi_u I_u + (1 - \lambda) \Phi_v I_v \\ &\leq \lambda \Phi_v I_u + (1 - \lambda) \Phi_v I_v = \Phi_v I_{\text{syn}}. \end{aligned} \quad (9)$$

Dividing throughout by $I_{\text{syn}} > 0$ yields $\Phi_u \leq \frac{O_{\text{syn}}}{I_{\text{syn}}} \leq \Phi_v$. Because $\Phi_u, \Phi_v \in [1 - \bar{\epsilon}, 1.0]$, it follows directly that $\Phi_{\text{flow}}(\mathbf{h}_{\text{syn}}) \in [1 - \bar{\epsilon}, 1.0]$. Hence, convex interpolation strictly preserves the bounded flow conservation property within the parent convex hull without violating monetary mass-balance bounds. $\square$

APPENDIX E

MASTER DATASET PROFILES & EXPERIMENTAL CONFIGURATIONS

A. Master Baseline Benchmark Matrix across All 14 Networks

Table S4 reports the comprehensive multi-dataset F1 and PR-AUC matrix across tabular, spatial, dynamic, and unsupervised architectures evaluated across 182 physical baseline runs.

B. 14-Dataset Hardware Profiling & Training Wall-Clock Compute

Table S5 details the compute hardware specifications, peak VRAM allocations, and per-epoch wall-clock training times on the Kaggle Cloud TPU / GPU compute nodes and local PC workstation testbed across all 14 financial transaction benchmarks [20]–[28].

C. Master Hyperparameter Configuration Grid

Table S6 provides the complete, exhaustive hyperparameter configuration across all six C-STGB modules.

D. Fine-Grained 12-D Invariant Feature Ablation

Table S7 reports leave-one-feature-out evaluations on Elliptic-v1.

E. Expanded 8-Pair Zero-Shot Cross-Domain Transfer Matrix

Table S8 reports zero-shot transfer performance across eight source-target dataset pairs with completely frozen neural backbones ($\eta = 0$).

Topological Invariant Isolation Test: To isolate the specific contribution of the 12-dimensional canonical invariant features ($\mathbf{z}_{\text{inv}}$) to cross-domain transfer, we evaluate an invariant ablation on the elliptic_v1 $\rightarrow$ elliptic_v2 transfer pair. Under full C-STGB with frozen backbone, the model achieves 82.40% zero-shot F1 (95.22% retention of supervised target performance). Removing the 12-D invariant projections collapses zero-shot F1 to 54.20% (a 28.20 pp drop), whereas vanilla HGT achieves only 21.40%. This establishes that the domain-invariant temporal and flow regularity signatures provide the primary inductive bridge enabling transfer across heterogeneous transaction networks.

F. GraphSMOTE Statistical Realism Metrics

Table S9 reports two-sample Kolmogorov-Smirnov statistics (D_{KS}) and Jensen-Shannon Divergence (JSD) between empirical illicit entities and GraphSMOTE-synthesized nodes on Elliptic-v1 ($N = 4,545$).

G. Conformal Temporal Drift Stress-Testing

Table S10 evaluates the calibrated conformal predictor across increasing out-of-time test horizons $\Delta t_{\text{drift}} \in [1, 4, 8, 12 \text{ weeks}]$ on continuous streaming transactions without parameter retraining, explicitly measuring the empirical total variation drift distance $d_{\text{TV}}(\mathcal{P}_t, \mathcal{P}_{\text{cal}})$ to validate the non-asymptotic coverage error bound.

H. Comparative Evaluation of Conformal Risk Control vs. Calibration Baselines

Table S11 compares Class-Conditional Conformal Risk Control against standard calibration and marginal conformal prediction baselines on elliptic_v1 at target significance $\alpha = 0.01$ (nominal coverage: 99.0%). Standard uncalibrated, temperature-scaled, and isotonic models lack finite-sample finite-set guarantees, yielding class-conditional coverage deficits on the minority class (90.10%–94.50%). Marginal conformal prediction satisfies aggregate coverage (99.05%) by over-covering majority benign transactions (99.98%) while severely under-covering minority illicit transactions (82.40%, a 16.60 pp violation). In contrast, C-STGB’s class-conditional formulation guarantees $\geq 99.0\%$ coverage for both classes simultaneously, isolating uncertainty into ambiguous sets ($\Gamma(X) = \{0, 1\}$) for 0.50% of transactions with zero empty sets.

TABLE S4

MASTER MULTI-BASELINE PERFORMANCE MATRIX ACROSS ALL 14 BENCHMARK NETWORKS: MACRO F1-SCORE (%) AND PR-AUC WITH VALIDATION-TUNED THRESHOLDS τ^* ACROSS TABULAR, SPATIAL, DYNAMIC, AND ANOMALY ARCHITECTURES ($\dagger$ DENOTES STATISTICALLY SIGNIFICANT SUPERIORITY OF C-STGB AGAINST DEEP GNN BASELINES AT $W = 105.0$, BENJAMINI-HOCHBERG ADJUSTED $p_{\text{Adj}} = 6.10 \times 10^{-4} < 0.001$ UNDER TWO-SIDED WILCOXON SIGNED-RANK TESTS ACROSS $n = 14$ PAIRED BENCHMARK NETWORKS AND 182 PHYSICAL BASELINE EVALUATION RUNS).

Dataset Identifier	XGBoost	CatBoost	LightGBM	Balanced RF	GCN	GraphSAGE	GAT	GIN	EvolveGCN	Logistic Reg	Autoencoder	C-STGB (Ours)
elliptic_v1	99.89/1.0000	99.78/1.0000	99.72/0.9999	95.35/0.9910	43.55/0.3482	49.07/0.4254	36.35/0.1959	57.32/0.5625	51.04/0.4243	58.51/0.4962	3.01/0.0352	99.44/1.0000[†]
elliptic_v2	99.81/0.9973	100.00/1.0000	100.00/1.0000	100.00/1.0000	0.00/0.0237	0.24/0.0230	0.00/0.0234	5.13/0.0266	0.00/0.0258	100.00/1.0000	5.03/0.0247	100.00/1.0000[†]
xblock_eth	96.91/0.9961	95.91/0.9925	97.05/0.9967	94.95/0.9857	6.53/0.0212	6.75/0.0206	6.45/0.0199	3.56/0.0137	6.62/0.0211	35.23/0.2941	8.08/0.0409	96.94/0.9915[†]
mtgox_leaked	74.39/0.8229	71.87/0.7983	73.93/0.8245	69.81/0.7459	20.30/0.2038	22.16/0.1801	22.47/0.1751	28.79/0.2411	22.12/0.0956	63.78/0.6854	0.09/0.2374	72.21/0.8306[†]
saml_d	93.74/0.9593	93.60/0.9448	92.85/0.9566	86.50/0.8633	1.69/0.0109	3.16/0.0071	3.20/0.0099	2.37/0.0078	3.81/0.0077	83.78/0.8104	28.07/0.1995	93.68/0.9574[†]
paysim1	18.90/0.1254	17.76/0.1061	1.32/0.3205	7.12/0.0351	3.50/0.0084	0.56/0.0016	0.46/0.0014	2.94/0.0057	3.98/0.0097	14.99/0.0865	8.06/0.0205	10.68/0.1173[†]
paysim_extended	99.72/0.9996	99.77/0.9996	99.80/0.9996	99.63/0.9994	96.22/0.9825	96.05/0.9799	96.64/0.9816	96.66/0.9818	92.64/0.7521	99.48/0.9992	0.52/0.2676	99.80/0.9993[†]
ibm_amsim_hi_small	36.66/0.3407	33.03/0.3100	36.62/0.3313	28.62/0.2772	2.46/0.0101	2.46/0.0080	1.03/0.0090	2.90/0.0112	2.31/0.0080	15.60/0.0850	0.92/0.0267	37.70/0.3550[†]
ibm_amsim_hi_medium	42.97/0.4513	41.53/0.4220	43.66/0.4563	36.72/0.3300	3.88/0.0135	3.95/0.0126	4.03/0.0164	3.95/0.0174	7.06/0.0350	28.84/0.2480	10.33/0.0510	42.85/0.4609[†]
ibm_amsim_li_small	15.31/0.1350	13.54/0.0802	12.86/0.1006	12.65/0.0600	0.84/0.0060	1.50/0.0057	0.61/0.0050	2.22/0.0105	1.18/0.0052	8.74/0.0425	3.03/0.0171	15.76/0.1485[†]
ibm_amsim_li_medium	23.19/0.1811	22.40/0.1718	23.55/0.1799	18.24/0.1270	3.41/0.0143	2.30/0.0073	4.05/0.0211	2.30/0.0083	4.29/0.0238	18.74/0.1134	5.02/0.0266	23.38/0.2077[†]
data_generator	100.00/1.0000	100.00/1.0000	99.99/1.0000	100.00/1.0000	99.74/0.9979	99.63/0.9956	99.98/1.0000	100.00/1.0000	62.72/0.6327	100.00/1.0000	82.72/0.9923	99.93/1.0000[†]
dgraphfin	98.23/0.9978	98.45/0.9989	98.59/0.9989	97.60/0.9873	2.60/0.0129	3.26/0.0161	2.71/0.0123	3.25/0.0189	2.59/0.0094	96.55/0.9431	0.42/0.0114	97.91/0.9979[†]
cc_transactions	51.15/0.5126	52.26/0.5092	52.76/0.5184	49.51/0.4967	48.16/0.3472	49.24/0.3555	48.62/0.3441	48.74/0.3883	4.79/0.0215	52.15/0.5085	48.24/0.4167	51.40/0.5213[†]
Macro-Average	67.92/0.6799	67.14/0.6666	66.62/0.6917	64.05/0.6356	23.78/0.2143	24.31/0.2170	23.33/0.2011	25.72/0.2353	18.94/0.1480	55.46/0.5223	14.54/0.1691	67.26/0.6848[†]

TABLE S5

COMPUTE HARDWARE SPECIFICATIONS AND TRAINING WALL-CLOCK RUNTIMES ACROSS ALL 14 BENCHMARKS.

Dataset Identifier	Entities ($ V $)	Edges ($ E $)	Peak VRAM	Sec / Epoch (C-STGB)
elliptic_v1	203,769	234,355	4.2 GB	12.4 s
elliptic_v2	122,279	186,400	3.1 GB	8.6 s
xblock_eth	2,150,000	2,800,000	11.2 GB	62.5 s
mtgox_leaked	145,000	210,000	2.8 GB	7.1 s
saml_d	980,000	1,450,000	8.4 GB	38.2 s
paysim1	6,362,620	6,362,620	12.4 GB	48.6 s
paysim_extended	1,048,575	1,048,575	6.9 GB	31.0 s
ibm_amsim_hi_small	100,000	180,000	2.2 GB	5.4 s
ibm_amsim_hi_medium	300,000	550,000	4.8 GB	16.8 s
ibm_amsim_li_small	100,000	180,000	2.2 GB	5.5 s
ibm_amsim_li_medium	300,000	550,000	4.9 GB	17.1 s
data_generator	100,000	250,000	2.6 GB	6.2 s
dgraphfin	3,700,550	1,604,218	2.3 GB	24.3 s
cc_transactions	284,807	284,807	3.4 GB	9.8 s

TABLE S6

COMPREHENSIVE MASTER HYPERPARAMETER SPECIFICATION FOR C-STGB.

Pipeline Module	Hyperparameter	Symbol	Value	Selection Method
1. Temporal Encoder	Burst Decay Prior	λ_{burst}	0.80	Grid Search ($\{0.5, 1.0\}$)
	Diurnal Decay Prior	λ_{diurnal}	0.05	24-Hour Periodicity
	Seasonal Decay Prior	$\lambda_{\text{seasonal}}$	0.01	90-Day Dormancy
	Minimum Weight Floor	w_{min}	0.05	Validation Split
2. Edge Trust Gate	Camouflage Gate Floor	δ_{floor}	0.10	Camouflage Sweep
	LeakyReLU Slope	α_{slope}	0.20	Standard Default
	Top-K Degree Cap	K	15	Latency/Memory Knee
3. GraphSMOTE	Typology Clusters	K_{clusters}	10	Silhouette Maximization
	Cosine Sim. Threshold	τ_{sim}	0.75	Cluster Cohesion
	Edge Reconstruct Thresh.	τ_{edge}	0.60	ROC-AUC Validation
4. Adaptive Fusion	Hidden Dimension	d	64	Memory Constraint
	Attention Heads	H	4	Multi-Head Sweep
	Dropout Rate	p	0.20	Regularization Sweep
5. Risk Policy	Task Loss Weight	γ_1, γ_2	0.50, 0.10	Pareto Multi-Task
	Cost Asymmetry Ratio	$C_{\text{FN}}/C_{\text{FP}}$	15.0	Regulatory Cost Model
	Target Risk Error	α_{risk}	0.01	Compliance Target
6. Conformal CRC	Significance Level	α	0.01 (99% Cov)	Fixed Policy Target
	Calibration Proportion	$ \mathcal{D}_{\text{cal}} / \mathcal{D} $	10%	Holdout Splitting

I. Calibration Sample Size Sensitivity and Minority Sample Thresholds

To evaluate calibration stability under extreme class imbalance, Table S12 reports empirical coverage and prediction set size across calibration split proportions $|\mathcal{D}_{\text{cal}}|/|\mathcal{D}| \in [1\%, 20\%]$ on elliptic_v1 at $\alpha = 0.01$. The theoretical condition $n_1 \geq \lceil (1 - \alpha)/\alpha \rceil = 99$ is satisfied across all splits $\geq 3\%$. Table S13 provides the exact empirical sample counts n_0 and n_1 in calibration splits across representative benchmark datasets.

TABLE S7

FINE-GRAINED CANONICAL INVARIANT FEATURE GROUP ABLATION ON ELLIPTIC-V1.

Invariant Configuration	Recall (%)	Precision (%)	F1-Score (%)
Full Invariant Representation (z_{inv})	98.88 $\pm$ 0.10	100.00 $\pm$ 0.00	99.44 $\pm$ 0.10
w/o 12-D Canonical Invariants (z_{inv})	96.80 $\pm$ 0.45	98.90 $\pm$ 0.35	97.84 $\pm$ 0.40 (−1.60 pp)
w/o Flow Conservation Ratio (Φ_{flow})	97.80 $\pm$ 0.16	96.10 $\pm$ 0.15	96.94 $\pm$ 0.14 (−2.50 pp)
w/o Hawkes Point Process Intensity ($\lambda_u(r)$)	98.05 $\pm$ 0.14	96.35 $\pm$ 0.12	97.19 $\pm$ 0.12 (−2.25 pp)
w/o Time-Causal PPR Taints ($s_{\text{fwd}}, s_{\text{bwd}}$)	97.20 $\pm$ 0.18	95.70 $\pm$ 0.15	96.44 $\pm$ 0.16 (−3.00 pp)
w/o 5-Moment Amount Statistics ($\mu_A \dots \nu_A$)	98.40 $\pm$ 0.12	98.10 $\pm$ 0.10	98.25 $\pm$ 0.11 (−1.19 pp)

TABLE S8

EXPANDED 8-PAIR ZERO-SHOT CROSS-DOMAIN TRANSFER MATRIX UNDER COMPLETELY FROZEN NEURAL BACKBONES ($\eta = 0$): C-STGB FROZEN REPORTS F1 / PR-AUC UNDER SOURCE THRESHOLD τ_{SRC}^* , WHILE C-STGB TARGET REPORTS F1 UNDER TARGET-CALIBRATED THRESHOLD τ_{TGT}^* .

Source Dataset	Target Dataset	HGT F1	XGB F1	C-STGB Frozen	C-STGB Target
elliptic_v1	elliptic_v2	21.40	64.20	82.40/0.8410	84.10
elliptic_v1	saml_d	14.20	58.10	76.80/0.7850	78.50
elliptic_v1	paysim_extended	16.80	61.40	79.20/0.8120	81.05
paysim_extended	ibm_amsim_hi	18.50	31.40	41.80/0.4320	43.50
paysim_extended	saml_d	22.10	56.80	74.50/0.7620	76.20
dgraphfin	elliptic_v1	19.40	62.80	80.60/0.8250	82.40
saml_d	ibm_amsim_hi	15.20	28.60	39.40/0.4080	41.10
xblock_eth	mtgox_leaked	28.50	69.20	85.10/0.8710	86.80
Macro-Average Zero-Shot Transfer		19.51	54.06	69.98/0.7170	71.71

J. Adversarial Camouflage Robustness across Perturbation Intensities

Table S14 reports model Macro F1-score retention on elliptic_v1 under increasing ratios of synthetic merchant camouflage edges ($\rho \in [0.0, 0.80]$). In addition to degree-matched merchant injection, we evaluate white-box gradient-guided edge additions (NETTACK-style). While vanilla HGT degrades by 50.85 pp (from 78.50% to 27.65%) and CARE-GNN degrades by 18.70 pp, C-STGB retains 89.63% F1 (a minimal −1.79 pp drop), confirming that learnable edge-trust gating successfully suppresses adversarial connections.

K. Cold-Start & Entity Degree Stratification Analysis

Table S15 reports model performance stratified across entity in/out degree regimes on Elliptic-v1 ($N = 203,769$). Crucially, this degree stratification dissects the division of labor between graph representation learning and tabular domain invariants: on isolated entities ($d = 1$), standalone GNNs suffer complete representation starvation (4.20% – 11.50% F1) due to neighborhood message depletion, whereas GBDT over 12-D

TABLE S9
STATISTICAL REALISM METRICS: EMPIRICAL ILLICIT VS. GRAPHSMOTE
SYNTHESIZED NODES ($N = 4, 545$).

Structural Metric Evaluated	KS Statistic (D_{KS})	p -value	JSD Metric
In/Out Degree Distribution	0.048	0.62 ($p > 0.50$)	0.018
Transaction Amount Distribution	0.035	0.84 ($p > 0.50$)	0.012
Inter-Arrival Temporal Dispersion ($\sigma_{\Delta t}$)	0.052	0.54 ($p > 0.50$)	0.021
Flow Conservation Ratio (Φ_{flow})	0.031	0.91 ($p > 0.50$)	0.009
Directed 3-Cycle Clustering Coeff.	0.042	0.76 ($p > 0.50$)	0.015

TABLE S10
CONFORMAL TEMPORAL DRIFT STRESS-TEST ON STREAMING PAYSIM OVER
INCREASING TIME HORIZONS (TARGET COVERAGE: $1 - \alpha = 99.0\%$, WITH
EMPIRICAL TOTAL VARIATION DRIFT d_{TV}).

Temporal Distance (Δt_{test})	Drift (d_{TV})	Benign Cov. (%)	Illicit Cov. (%)	Tier 2 Queue (%)	F1-Score (%)
1 Week Post-Calibration	0.012	99.12 $\pm$ 0.08	99.04 $\pm$ 0.12	0.42 $\pm$ 0.04	92.40 $\pm$ 0.60
4 Weeks Post-Calibration	0.038	99.06 $\pm$ 0.10	98.92 $\pm$ 0.15	0.58 $\pm$ 0.06	91.80 $\pm$ 0.65
8 Weeks Post-Calibration	0.074	98.85 $\pm$ 0.14	98.40 $\pm$ 0.22	0.94 $\pm$ 0.09	90.15 $\pm$ 0.72
12 Weeks Post-Calibration	0.126	98.40 $\pm$ 0.18	97.60 $\pm$ 0.30	1.85 $\pm$ 0.15	88.20 $\pm$ 0.85

invariants alone delivers 84.20% F1. The C-STGB gating network adaptively assigns $\bar{\alpha}_u = 0.08$, smoothly leveraging the invariant foundation while gleaning subtle 1-hop contextual signals, elevating performance to 89.60% F1. As relational degree increases ($d \geq 20$), the gating parameter shifts to $\bar{\alpha}_u = 0.91$, fully unlocking high-order graph relational patterns (93.80% F1).

L. Disentangled Production Streaming Latency Profiling

Table S16 reports the itemized computational budget across individual pipeline stages on continuous streaming financial transactions, disentangling the lightweight Fast-Path (0.45 ms) from the comprehensive Full-Path (2.10 ms, including GNNExplainer factual subgraph extraction for Tier 1 alerts).

M. 5-Order-of-Magnitude Systems Scalability Stress-Testing

Table S17 reports latency and memory scaling across graph sizes ranging from 10^3 to 5×10^6 nodes under degree capping ($K \leq 15$) on benchmark hardware (AMD 5900X, RTX 3080). Per-node memory overhead remains bounded at ≈ 8.63 – 8.97 KB, and batch inference latency scales sublinearly from 3.32 ms to 3.48 ms, consistently satisfying the 35 ms industrial banking SLA.

N. Asymmetric Regulatory Cost-Benefit Analysis

Table S18 quantifies the operational loss $\mathcal{L}_{\text{ops}} = C_{\text{FN}} \cdot \text{FN} + C_{\text{FP}} \cdot \text{FP}$ across increasing regulatory penalty ratios $C_{\text{FN}}/C_{\text{FP}} \in [1, 50]$ on Elliptic-v1, demonstrating that Dynamic Bayes thresholding τ^* achieves up to 78.4% cost reduction over uncalibrated thresholds ($\tau = 0.50$).

O. Hyperparameter Sensitivity & Stability Analysis

Table S19 evaluates the stability of C-STGB on Elliptic-v1 across sweeps of key architectural hyperparameters.

APPENDIX F REPRODUCIBILITY CHECKLIST

To ensure full compliance with IEEE Signal Processing Society and TIFS deep learning reproducibility standards, Table S20 provides the complete software, hardware, environment, and seed configuration specification.

APPENDIX G

14.DATASET INVARIANT-AUGMENTED GNN CROSS-EVALUATION MATRIX

Table S21 reports the complete multi-dataset performance breakdown across all 14 benchmark networks for spatial and dynamic GNN architectures (GCN, GraphSAGE, and EvolveGCN) trained with canonical node attributes concatenated with the full 12-dimensional domain invariant vector $[\mathbf{x}_u \parallel \mathbf{z}_{\text{inv}}(u, t)]$.

As detailed in Section IV-C of the main manuscript, invariant concatenation yields only marginal improvements (+0.55 pp to +1.11 pp in macro F1-score) over canonical feature inputs. This confirms the mechanistic finding that isotropic neighborhood message aggregation linearly averages localized scalar invariant signatures (such as flow conservation $\Phi_{\text{flow}} \approx 1.0$ and Hawkes arrival intensity λ_u) across high-degree benign hubs, diluting the anomaly signal into the majority background mean. Crucially, C-STGB maintains an unbroken 14/0/0 win record against all three invariant-augmented GNN architectures across all 14 empirical benchmarks ($W = 105.0$, $W_- = 0.0$, $p = 1.22 \times 10^{-4}$, $p_{\text{adj}} = 6.10 \times 10^{-4} < 0.001$), confirming that the statistical significance and architectural superiority of C-STGB are strictly invariant to baseline feature representation.

APPENDIX H

TOPOLOGICAL EXPLANATION FIDELITY, FACTUAL SPARSITY, AND COUNTERFACTUAL VALIDATION

To substantiate model explainability under regulatory scrutiny (such as Federal Reserve SR 26-2 [16] and EU AI Act Article 13 transparency mandates), this section formalizes the quantitative attribution fidelity, explanation sparsity, and counterfactual validation of C-STGB’s GNNExplainer topological extraction pipeline, evaluated across $N = 500$ confirmed illicit transaction clusters from `elliptic_v1`, `paysim_extended`, and `saml_d` (Table S22).

A. Mathematical Definitions of Attribution Metrics

- 1) **Factual Subgraph Fidelity (Fid^+)**: Quantifies whether the isolated causal subgraph $\mathcal{G}_s$ captures the sufficient and necessary relational topology triggering the anomaly classification:

$$\text{Fid}^+ = \frac{1}{N} \sum_{i=1}^N (\hat{p}(y_i = 1 \mid \mathcal{G}_i) - \hat{p}(y_i = 1 \mid \mathcal{G}_i \setminus \mathcal{G}_{s,i})) \quad (10)$$

A high Fid^+ score (0.942) demonstrates that removing the salient subgraph collapses illicit confidence to baseline benign levels.

- 2) **Counterfactual Infidelity (Fid^-)**: Measures whether background chaff edges outside the explanation mask contain residual illicit information:

$$\text{Fid}^- = \frac{1}{N} \sum_{i=1}^N (\hat{p}(y_i = 1 \mid \mathcal{G}_i) - \hat{p}(y_i = 1 \mid \mathcal{G}_{s,i})) \quad (11)$$

A near-zero Fid^- score (0.018) confirms that the isolated subgraph alone preserves full model prediction accuracy without uncaptured structural residuals.

TABLE S11

COMPARATIVE EVALUATION OF CONFORMAL RISK CONTROL AGAINST CALIBRATION AND CONFORMAL BASELINES ON ELLIPTIC-V1 (TARGET SIGNIFICANCE $\alpha = 0.01$, NOMINAL COVERAGE $1 - \alpha = 99.0\%$).

Calibration / Conformal Method	Overall Cov. (%)	Benign Cov. (%)	Illicit Cov. (%)	Illicit Cov. Error	Mean Set Size $\mathbb{E}[\Gamma]$	Empty Sets (%)
Uncalibrated Cutoff ($\tau = 0.50$)	91.42	99.85	90.10	-8.90 pp	1.000	0.00%
Temperature Scaling [29]	94.20	99.70	92.80	-6.20 pp	1.000	0.00%
Isotonic Regression [30]	96.10	99.65	94.50	-4.50 pp	1.000	0.00%
Marginal Conformal Prediction (Split CP) [14]	99.05	99.98	92.40	-16.60 pp	1.001	0.00%
Class-Conditional CRC (Ours, Guarantee A)	99.05	99.05	99.12	+0.12 pp	1.005	0.00%
Streaming Adaptive Conformal Inference (Ours, Guarantee C)	99.04	99.06	99.08	+0.08 pp	1.006	0.00%

TABLE S12

CALIBRATION SPLIT SIZE SENSITIVITY ON ELLIPTIC-V1 (TARGET $\alpha = 0.01$, NOMINAL COVERAGE 99.0%).

Calibration Split (%)	Empirical Cov. (%)	Guarantee Satisfied	Mean Set Size $\mathbb{E}[\Gamma]$	Ambiguous Rate (%)
1%	99.02	✓	1.0070	0.70%
3%	99.06	✓	1.0049	0.49%
5%	99.08	✓	1.0042	0.42%
10% (Default)	99.09	✓	1.0036	0.36%
15%	99.09	✓	1.0033	0.33%
20%	99.10	✓	1.0031	0.31%

TABLE S13

EMPIRICAL CALIBRATION SAMPLE COUNTS (n_0, n_1) ACROSS KEY BENCHMARK NETWORKS (10% CALIBRATION SPLIT).

Dataset Identifier	Total Nodes	Benign (n_0)	Illicit (n_1)	Condition $n_1 \geq 100$
elliptic_v1	203,769	15,700	454	✓(4.5× margin)
elliptic_v2	122,279	10,200	203	✓(2.0× margin)
saml_d	980,000	97,200	800	✓(8.0× margin)
paysim_extended	1,048,575	104,700	157	✓(1.6× margin)
ibm_amsim_hi_small	100,000	9,950	50	Historical pooling required ($n_1 \geq 100$)

- 3) **Explanation Sparsity:** Measures the fraction of the 2-hop computation graph pruned as commercial or benign noise:

$$\text{Sparsity} = 1 - \frac{|\mathcal{E}_s|}{|\mathcal{E}_c|} \quad (12)$$

where $|\mathcal{E}_c|$ is the computational ego-graph size and $|\mathcal{E}_s|$ is the pruned causal conduit (86.5% average compression).

- 4) **Counterfactual Validity:** The empirical rate at which normalizing attributed transaction amounts or diluting burst frequency successfully transitions borderline or high-risk accounts from Tier 3 to Tier 1 ($\Gamma(X) = \{0\}$), establishing actionable operational recourse.

B. Pilot Human Decision-Support Compliance Evaluation

To evaluate the operational utility of the auxiliary multi-agent SAR drafting prototype (Section VI-C of the main manuscript), Table S23 reports findings from a blinded pilot human evaluation conducted with two certified Bank Secrecy Act (BSA) compliance examiners across $N = 200$ audit dossiers (100 manual investigations and 100 C-STGB prototype-assisted dossiers). Evaluated against statutory FinCEN Form 111 standards, the prototype achieved 100.0% entity and dollar-amount factual grounding, 98.5% typology identification concordance ($\kappa = 1.0$), and 100.0% statutory citation correctness. Crucially, the prototype reduced average investigation and draft compilation time from 45.0 minutes to 28.4 seconds per case (95.6× speedup), confirming substantial operational decision-support utility under mandatory human officer sign-off.

TABLE S14

ADVERSARIAL CAMOUFLAGE STRESS-TEST ON ELLIPTIC-V1: MACRO F1-SCORE (%) UNDER INCREASING PERTURBATION RATIO ρ .

Camouflage Ratio (ρ)	C-STGB (Ours)	CARE-GNN	Vanilla HGT	GCN
0.0 (Clean)	91.42	82.00	78.50	18.73
0.1	91.39	80.75	73.34	17.23
0.2	91.31	78.91	67.43	15.73
0.4	90.97	74.40	54.78	12.73
0.6	90.41	69.13	41.44	9.73
0.8	89.63	63.30	27.65	6.73

TABLE S15

PERFORMANCE STRATIFICATION ACROSS ENTITY DEGREE REGIMES ON ELLIPTIC-V1.

Degree Regime [$d(u)$]	Prevalence	GCN F1	Vanilla HGT F1	CARE-GNN F1	XGBoost F1	C-STGB F1 ($\bar{a}_u$)
Isolated Mules [$d = 1$]	38.4%	4.20%	11.50%	24.80%	84.20%	89.60% ($\bar{a}_u = 0.08$)
Low-Degree Conduits [$2 \leq d \leq 5$]	42.1%	16.80%	42.10%	58.40%	87.10%	90.80% ($\bar{a}_u = 0.44$)
Medium Meshes [$6 \leq d \leq 19$]	14.5%	32.40%	68.40%	74.20%	89.40%	92.50% ($\bar{a}_u = 0.82$)
High-Degree Hubs [$d \geq 20$]	5.0%	41.50%	74.20%	78.50%	91.20%	93.80% ($\bar{a}_u = 0.91$)

REFERENCES

- [1] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 785–794.
- [2] L. Prokhorenkova, G. Gusev, A. Vorobev, A. V. Dorogush, and A. Gulin, “CatBoost: Unbiased boosting with categorical features,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 31, 2018.
- [3] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, “LightGBM: A highly efficient gradient boosting decision tree,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017, pp. 3146–3154.
- [4] T. N. Kipf and M. Welling, “Semi-supervised classification with graph convolutional networks,” in *International Conference on Learning Representations (ICLR)*, 2017.
- [5] P. Veličković, G. Cucurull, A. Casanova, A. Romero, P. Liò, and Y. Bengio, “Graph attention networks,” in *International Conference on Learning Representations (ICLR)*, 2018.
- [6] K. Xu, W. Hu, J. Leskovec, and S. Jegelka, “How powerful are graph neural networks?” in *International Conference on Learning Representations (ICLR)*, 2019.
- [7] A. Pareja, G. Domeniconi, J. Chen, T. Teng, T. Wang, K. Chai, M. Magdon-Ismael, and C. Shen, “EvolveGCN: Evolving graph convolutional networks for dynamic graphs,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 34, no. 04, 2020, pp. 5363–5370.
- [8] E. Rossi, B. Zhou, F. Monti, F. Frasca, U. Alon, and M. M. Bronstein, “Temporal graph networks for deep learning on dynamic graphs,” in *NeurIPS Workshop on Graph Representation Learning and Beyond*, 2020.
- [9] D. Xu, C. Ruan, E. Korpeoglu, S. Kumar, and K. Achan, “Inductive representation learning on temporal graphs,” in *International Conference on Learning Representations (ICLR)*, 2020.
- [10] Y. Dou, Z. Liu, L. Sun, Y. Deng, H. Peng, and P. S. Yu, “Enhancing graph neural network-based fraud detectors against camouflaged fraudsters,” in *Proceedings of the 29th ACM International Conference on Information & Knowledge Management (CIKM)*, 2020, pp. 315–324.
- [11] T. Zhao, X. Zhang, and S. Wang, “GraphSMOTE: Imbalanced node classification on graphs with graph neural networks,” in *Proceedings of the 14th ACM International Conference on Web Search and Data Mining (WSDM)*, 2021, pp. 833–841.

Comprehensive Small-Multiples Precision-Recall Curves across ALL 14 Empirical Datasets




Fig. S2. Comprehensive small-multiples Precision-Recall curves across all 14 empirical financial transaction benchmark networks (plus Global Macro-Average summary), showing PR-AUC scores across varying class skews ($\pi \in [0.05\%, 2.23\%]$).

TABLE S16

FINE-GRAINED LATENCY PROFILING PER TRANSACTION EVENT (CPU: AMD 5900X, GPU: RTX 3080).

Processing Stage	Mean Latency	P95 Latency	P99 Latency
1. DuckDB Ingestion & 12-D Invariant Extraction	0.11 ms	0.16 ms	0.22 ms
2. Dynamic Top- K ($K \leq 15$) Subgraph Retrieval	0.18 ms	0.26 ms	0.35 ms
3. Temporal GNN & Evidence-Adaptive Forward Pass	0.12 ms	0.17 ms	0.21 ms
4. Conformal CRC Calibration & 3-Tier Routing	0.04 ms	0.05 ms	0.07 ms
Total End-to-End In-Flight Fast-Path	0.45 ms	0.64 ms	0.82 ms
5. Post-Hoc GNNExplainer Subgraph Extraction (Tier 1 Alerts)	1.65 ms	2.10 ms	2.85 ms
Total End-to-End Full-Path (with Forensic Attribution)	2.10 ms	2.74 ms	3.67 ms
Batch-64 Amortized Throughput (GPU)	0.054 ms	0.072 ms	0.095 ms

TABLE S17

5-ORDER-OF-MAGNITUDE SYSTEMS SCALABILITY STRESS-TEST (RTX 3080, DEGREE CAP $K = 15$).

Node Count ($ V $)	Batch Latency (ms)	Memory (KB / Node)	SLA Satisfied (≤ 35 ms)
1,000	3.32	8.63	✓
10,000	3.35	8.70	✓
100,000	3.40	8.80	✓
1,000,000	3.45	8.90	✓
5,000,000	3.48	8.97	✓

TABLE S18

EXPECTED OPERATIONAL COMPLIANCE COST UNDER ASYMMETRIC REGULATORY RATIOS (NORMALIZED TO $\tau = 0.50$).

Cost Asymmetry (C_{FN}/C_{FP})	GCN ($\tau = 0.50$)	XGBoost (τ^*)	C-STGB ($\tau = 0.50$)	C-STGB (τ^*)
Symmetric Loss (1 : 1)	1.000	0.284	0.215	0.142 (−85.8%)
Standard Compliance (5 : 1)	1.000	0.342	0.280	0.178 (−82.2%)
Strict BSA/AML (15 : 1, Default)	1.000	0.395	0.345	0.216 (−78.4%)
Severe Regulatory Sanction (30 : 1)	1.000	0.448	0.398	0.245 (−75.5%)
Critical Anti-Terrorism (50 : 1)	1.000	0.490	0.435	0.272 (−72.8%)

- [12] F. Johannessen and M. Jullum, “Finding money launderers using heterogeneous graph neural networks,” *The Journal of Finance and Data Science*, vol. 11, p. 100175, 2025.
- [13] Y. Wang, K. Ding, C. Zhou, J. Liu, and S. Parthasarathy, “Dynamic graph anomaly detection with multi-scale temporal graph neural networks,” in *Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*, 2024, pp. 3210–3221.
- [14] V. Vovk, A. Gammerman, and G. Shafer, *Algorithmic Learning in a Random World*. Springer Science & Business Media, 2005.
- [15] T. Ding, A. N. Angelopoulos, S. Bates, M. I. Jordan, and R. J. Tibshirani,

“Class-conditional conformal prediction with many classes,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 36, 2023, pp.




Fig. S3. Adversarial Camouflage Topological Denoising: (a) Injected 50-edge noisy merchant topology where an illicit laundering cycle deliberately injects links to high-degree commercial hubs. (b) C-STGB edge-gating module dynamically suppresses 65.9% of spurious links ($g_{ij} < 0.10$), cleanly isolating the 4-hop wash ring ($g_{ij} \geq 0.92$) and preventing feature over-smoothing.

Multi-Agent Autonomous Forensic Swarm & Fed SR 11-7 Cryptographic Audit Pipeline




Fig. S4. Auxiliary Multi-Agent Assisted Decision Support and Model-Governance Cryptographic Audit Pipeline (Aligned with Federal Reserve SR 26-2 [16] and EU AI Act governance standards): (1) Investigator Agent extracts model-based attribution subgraphs; (2) SAR Drafter Agent compiles structured research XML narratives under FinCEN Report 111 guidelines; (3) Compliance Auditor checks BSA/AML regulatory constraints; and (4) Model Governance Logger generates tamper-evident SHA-256 Merkle chain receipts.

TABLE S19
HYPERPARAMETER SENSITIVITY EVALUATION ON ELLIPTIC-V1.

Hyperparameter Swept	Tested Range	Optimal Value	F1-Score Stability Range	Operational Impact
Edge-Trust Floor (ϕ_{trust})	[0.01, 0.30]	0.10	99.10% – 99.44%	Broad stable plateau in [0.05, 0.15]
Top- K Degree Cap (K)	[5, 35]	15	98.80% – 99.44%	Latency/Recall Pareto knee at $K = 15$
Cluster Count (K_{clusters})	[3, 20]	10	99.05% – 99.44%	Robust across $K \geq 8$ typologies
Significance Level (α)	[0.001, 0.05]	0.01	Coverage: 99.9% → 95.1%	Tier 2 queue shrinks from 1.20% → 0.18%

4461–4475.

- [16] Board of Governors of the Federal Reserve System, “Supervisory letter sr 26-2: Interagency guidance on model risk management for advanced machine learning and artificial intelligence technologies,” Federal Reserve Board, Office of the Comptroller of the Currency, and FDIC,

Washington, DC, Tech. Rep., 2026, issued April 17, 2026; supersedes SR 11-7 and SR 21-8.

- [17] R. J. Tibshirani, E. J. Candès, A. Ramdas, and R. Foygel Barber, “Conformal prediction under covariate shift,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.
- [18] R. F. Barber, E. J. Candès, A. Ramdas, and R. J. Tibshirani, “Conformal prediction beyond exchangeability,” *The Annals of Statistics*, vol. 51, no. 2, pp. 816–845, 2023.
- [19] I. Gibbs and E. J. Candès, “Adaptive conformal inference under distribution shift,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, 2021, pp. 1660–1672.
- [20] M. Weber, G. Domeniconi, J. Chen, K. D. Weidele, C. Bellei, T. Robinson, and C. Shen, “Anti-Money Laundering in Bitcoin: Experimenting with graph convolutional networks for financial forensics,” in *Proceed-*




Fig. S5. Canonical 12-Dimensional Spatio-Temporal Domain Invariant Representations: Detailed schematic illustrating the extraction of flow conservation ratios Φ_{flow} , degree asymmetry Ψ_{asym} , temporally restricted PageRank taint $s_{\text{fwd}}/s_{\text{bwd}}$, continuous Hawkes process arrival intensity $\lambda_u(t)$, and higher-order statistical moments across directed transaction subgraphs.




Fig. S6. Evidence-Adaptive Cold-Start Resolution and Tri-Band Spectral Response: (a) Dynamic transition of the adaptive gating parameter $\alpha_u \in [0, 1]$ as a function of entity degree $\deg(u)$, demonstrating reliance on 12-D canonical physical invariants for disconnected cold-start accounts ($\deg(u) = 0$). (b) Spectral frequency decomposition of the Tri-Band continuous harmonic attention bank across burst, diurnal, and seasonal transaction velocity regimes.

TABLE S20

DEEP LEARNING REPRODUCIBILITY CHECKLIST & EXECUTION ENVIRONMENT.

Reproducibility Dimension	Specification / Value
Code Repository	https://anonymous.4open.science/r/Intelligent-AML-Review
Primary Programming Language	Python 3.10.12
Deep Learning Framework	PyTorch 2.3.0+cu121, PyTorch Geometric (PyG) 2.5.2
Graph Kernel Engine	DGL 2.2.1, NetworkX 3.2.1
Relational Data Engine	DuckDB 0.10.2, Apache Arrow 16.0.0
Scientific Stack	NumPy 1.26.4, SciPy 1.13.0, Scikit-learn 1.4.2
Conformal Prediction Library	MAPIE 0.8.0, Custom Class-Conditional CRC Engine
Hardware Environment	Kaggle Cloud TPU v3-8 / NVIDIA Tesla T4 GPU (16GB)
Workstation Environment	AMD Ryzen 9 5900X, 64GB DDR4 RAM, RTX 3080 (10GB)
Operating System	Ubuntu 22.04 LTS / Windows 11 Enterprise (x86_64)
Random Seeds Evaluated	$seeds \in \{42, 101, 2024, 7, 999\}$
Master Training Command	<code>python scripts/train_unified_pipeline.py --dataset all</code>
Benchmark Evaluation Script	<code>python scripts/benchmark_baselines.py --runs 5</code>
Figure Generator Script	<code>python scripts/generate_all_publication_figures.py</code>

ings of the ACM SIGKDD Workshop on Anomaly Detection in Finance, 2019, pp. 1–8.

- [21] C. Bellei, M. Xu, R. Phillips, T. Robinson, M. Weber, T. Kaler, C. E. Leiserson, Arvind, and J. Chen, “The shape of money laundering: Subgraph representation learning on the blockchain with the Elliptic2 dataset,” *arXiv preprint arXiv:2404.19109*, 2024.
- [22] X. Huang, Y. Shen, Y. Huang, Z. Zhang, X. Zhou, X. Zhang, Y. Du, and B. Hu, “DGraph: A large-scale financial dataset for graph anomaly detection,” in *Advances in Neural Information Processing Systems (NeurIPS) Track on Datasets and Benchmarks*, vol. 35, 2022, pp. 22 765–22 777.
- [23] P. Zheng, H.-N. Dai, J. Wu, and Z. Zheng, “XBlock-ETH: Extracting and analyzing blockchain data from Ethereum,” *IEEE Open Journal of the Computer Society*, vol. 1, pp. 95–106, 2020.
- [24] N. Gandal, J. T. Hamrick, T. Moore, and T. Oberman, “Price manipula-

tion in the Bitcoin ecosystem,” *Journal of Monetary Economics*, vol. 95, pp. 86–96, 2018.

- [25] J. Lorenz, M. I. Silva, D. Aparício, J. T. Ascensão, and P. Bizarro, “Machine learning applied to anomaly detection in financial transactions: A review and an application to a multi-bank dataset,” *arXiv preprint arXiv:2010.05779*, 2020.
- [26] E. A. Lopez-Rojas and S. Axelsson, “PaySim: A financial mobile money simulator for fraud detection,” in *Proceedings of the 28th European Modeling and Simulation Symposium (EMSS)*, 2016, pp. 249–255.
- [27] T. Suzumura, Y. Zhou, N. Baracaldo, S. Guan, K. Hou, R. Houlihan, M. Stamatatos, P. Sapkota, and C. Zhang, “Towards online Anti-Money Laundering detection on distributed dynamic graphs,” in *Proceedings of the IEEE International Conference on Big Data (Big Data)*, 2019, pp. 2350–2359.
- [28] A. Dal Pozzolo, O. Caelen, R. A. Johnson, and G. Bontempi, “Calibrating probability with undersampling for unbalanced classification,” in *IEEE Symposium Series on Computational Intelligence (SSCI)*, 2015, pp. 159–166.
- [29] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, “On calibration of modern neural networks,” in *Proceedings of the 34th International Conference on Machine Learning (ICML)*, 2017, pp. 1321–1330.
- [30] B. Zadrozny and C. Elkan, “Transforming classifier scores into accurate multiclass probability estimates,” in *Proceedings of the 8th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2002, pp. 694–699.

TABLE S21

MASTER INVARIANT-AUGMENTED GNN CROSS-EVALUATION MATRIX ACROSS ALL 14 FINANCIAL NETWORKS: MACRO F1-SCORE (%) AND PR-AUC WITH VALIDATION-TUNED THRESHOLDS τ^* COMPARING CANONICAL INPUTS $[\mathbf{x}_{it}]$ VS. INVARIANT-AUGMENTED REPRESENTATIONS $[\mathbf{x}_{it} \parallel \mathbf{z}_{inv}]$ AGAINST PROPOSED C-STGB († DENOTES STATISTICALLY SIGNIFICANT SUPERIORITY OF C-STGB AT $W = 105.0$, $p_{adj} < 0.001$ UNDER TWO-SIDED WILCOXON SIGNED-RANK TESTS WITH 14 WINS AND 0 LOSSES).

Dataset Identifier	GCN (Canon.)	GCN (+ $\mathbf{z}_{inv}$)	GraphSAGE (Canon.)	GraphSAGE (+ $\mathbf{z}_{inv}$)	EvolveGCN (Canon.)	EvolveGCN (+ $\mathbf{z}_{inv}$)	C-STGB (Proposed)
elliptic_v1	43.55/0.3482	44.82/0.3590	49.07/0.4254	50.45/0.4380	51.04/0.4243	51.85/0.4312	99.44/1.0000 [†]
elliptic_v2	0.00/0.0237	0.45/0.0245	0.24/0.0230	0.88/0.0242	0.00/0.0258	0.22/0.0261	100.00/1.0000 [†]
xblock_eth	6.53/0.0212	7.42/0.0225	6.75/0.0206	7.81/0.0221	6.62/0.0211	7.15/0.0219	96.94/0.9915 [†]
mtgox_leaked	20.30/0.2038	21.45/0.2110	22.16/0.1801	23.35/0.1895	22.12/0.0956	22.84/0.1020	72.21/0.8306 [†]
saml_d	1.69/0.0109	2.45/0.0118	3.16/0.0071	3.92/0.0082	3.81/0.0077	4.30/0.0084	93.68/0.9574 [†]
paysim1	3.50/0.0084	4.15/0.0092	0.56/0.0016	1.12/0.0022	3.98/0.0097	4.45/0.0105	10.68/0.1173 [†]
paysim_extended	96.22/0.9825	96.85/0.9840	96.05/0.9799	96.72/0.9815	92.64/0.7521	93.20/0.7610	99.80/0.9993 [†]
ibm_amlsim_hi_small	2.46/0.0101	3.20/0.0112	2.46/0.0080	3.15/0.0091	2.31/0.0080	2.85/0.0088	37.70/0.3550 [†]
ibm_amlsim_hi_medium	3.88/0.0135	4.65/0.0148	3.95/0.0126	4.70/0.0139	7.06/0.0350	7.62/0.0365	42.85/0.4609 [†]
ibm_amlsim_li_small	0.84/0.0060	1.45/0.0071	1.50/0.0057	2.10/0.0066	1.18/0.0052	1.65/0.0059	15.76/0.1485 [†]
ibm_amlsim_li_medium	3.41/0.0143	4.10/0.0155	2.30/0.0073	2.95/0.0082	4.29/0.0238	4.82/0.0250	23.38/0.2077 [†]
data_generator	99.74/0.9979	99.82/0.9982	99.63/0.9956	99.71/0.9962	62.72/0.6327	63.45/0.6401	99.93/1.0000 [†]
dgraphfin	2.60/0.0129	3.25/0.0141	3.26/0.0161	3.90/0.0175	2.59/0.0094	3.10/0.0102	97.91/0.9979 [†]
cc_transactions	48.16/0.3472	48.95/0.3540	49.24/0.3555	50.05/0.3620	4.79/0.0215	5.35/0.0232	51.40/0.5213 [†]
Macro-Average	23.78/0.2143	24.85/0.2215	24.31/0.2170	25.42/0.2248	18.94/0.1480	19.49/0.1528	67.26/0.6848 [†]

TABLE S22

QUANTITATIVE ATTRIBUTION FIDELITY, SPARSITY, AND COUNTERFACTUAL VALIDATION OF GNNEXPLAINER SUBGRAPH MASKS ACROSS CONFIRMED ILLICIT CLUSTERS ($N = 500$ CLUSTERS FROM ELLIPTIC-V1, PAYSIM EXTENDED, AND SAML-D).

Metric Evaluated	Regulatory Target	Empirical Result	Forensic Standard
Factual Motif Precision	$\geq 90.0\%$	96.40 $\pm$ 0.85%	Isolates true smurfing loop
Factual Motif Recall	$\geq 90.0\%$	94.80 $\pm$ 0.92%	Captures full layering conduit
Explanation Sparsity	$\geq 80.0\%$	86.50 $\pm$ 1.20%	Prunes benign merchant noise
Counterfactual Validity	$\geq 95.0\%$	98.20 $\pm$ 0.45%	Validates operational recourse
Composite Fidelity (Fid*)	≥ 0.90	0.942 $\pm$ 0.015	High explanation faithfulness
Infidelity Bias (Fid ⁻)	≤ 0.05	0.018 $\pm$ 0.006	Minimal background residual
Attribution Concordance (κ)	≥ 0.95	0.980 $\pm$ 0.010	Inter-examiner audit agreement

TABLE S23

PILOT HUMAN DECISION-SUPPORT COMPLIANCE EVALUATION ACROSS $N = 200$ CASES WITH 2 BSA EXAMINERS.

Audit Dimension Evaluated	Target Forensic Threshold	Empirical Score
Entity Identifier Accuracy	100.0% (Zero Hallucination)	100.0% (200/200)
Transaction Amount Factual Grounding	100.0% (Exact Matching)	100.0% (200/200)
Laundering Typology Correctness	$\geq 95.0\%$	98.5% ($\kappa = 1.00$)
Statutory Citation Correctness	100.0% (Valid CFR / USC)	100.0% (200/200)
Average Investigation Time (Manual)	Baseline Review	45.0 min / case
Average Investigation Time (C-STGB)	Real-Time Decision Support	28.4 s / case
Operational Investigation Speedup	$\geq 10\times$	95.6$\times$

14-Dataset Multi-Criteria Radar Benchmark




Fig. S7. Multi-Dimensional Empirical Radar Comparisons across the 14 Financial Transaction Benchmarks: Visualizing relative performance across Precision, Recall, Macro F1, PR-AUC, Adversarial Robustness, and Sub-Millisecond Inference Latency against competitive baselines.




Fig. S8. Typology-Clustered Latent GraphSMOTE Virtual Node and Parametric Bilinear Edge Generation Pipeline: (a) Partitioning minority illicit entities into latent typology clusters C_k via Spherical k -Means; (b) Linear interpolation synthesis along cluster manifolds; and (c) Bilinear candidate link scoring satisfying bounded mass-balance flow conservation $\Phi_{\text{flow}} \in [\min, \max]$ (Proposition S2) and topological zero-leakage isolation from evaluation splits.